

Concurrency

Edward Z. Yang

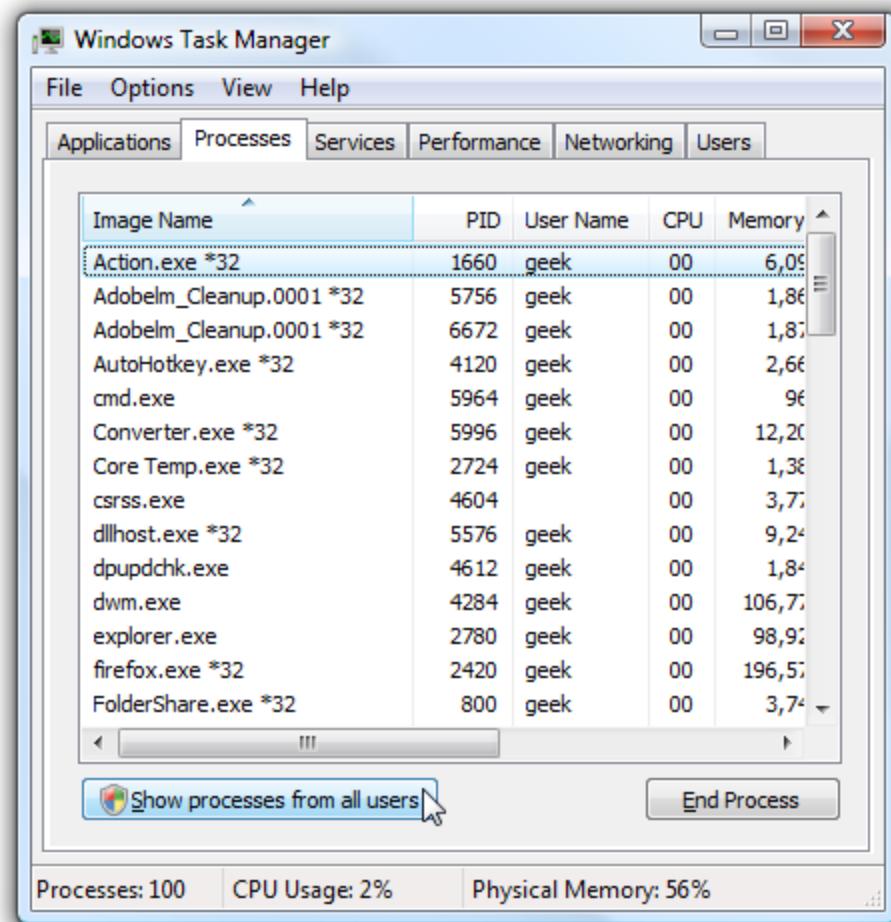
What is concurrency?

"It's when you have multiple threads"

"Things execute at the same time"

Why concurrency?

Concurrency is ubiquitous



Concurrency is ubiquitous



How can programming languages make concurrent programming easier?

What abstractions are most effective?



Message Passing
Actors CSP

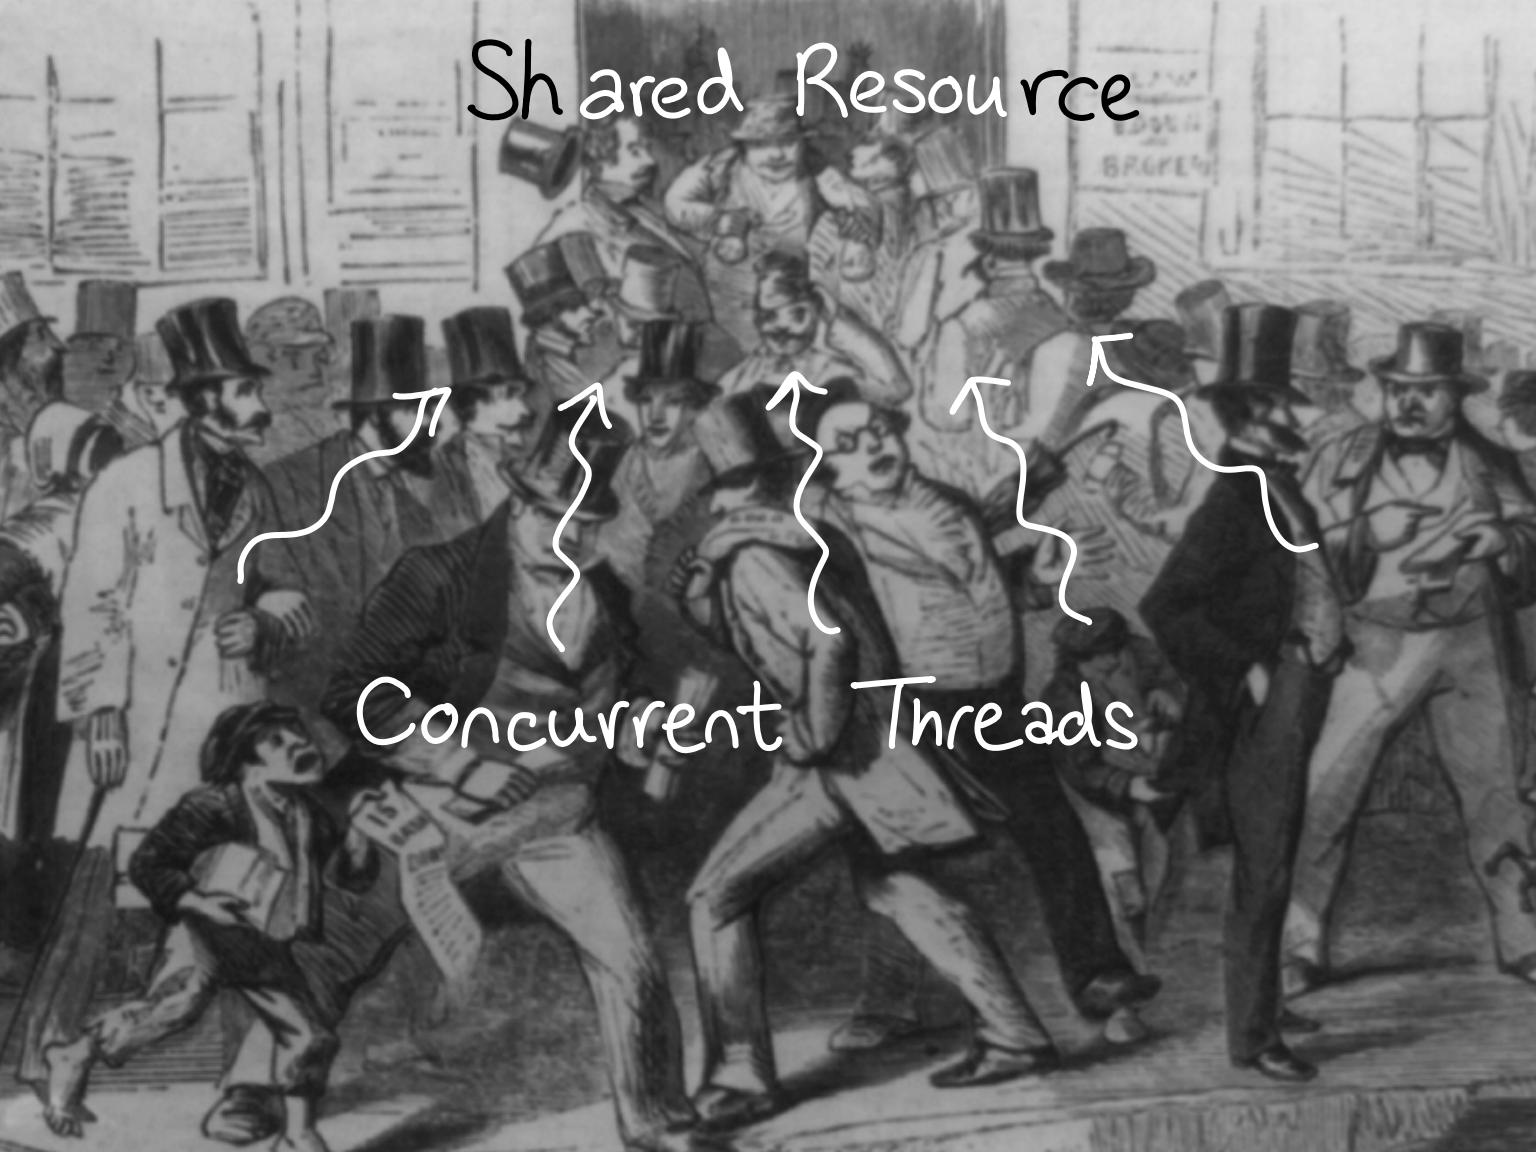
Transactional
Memory

Monitors
(Java)

Deterministic
Parallelism

Locks

Memory Model



Shared Resource

Concurrent Threads

Memory = Shared Resource

balance := 100

Thread #1

```
if (balance ≥ 100) {  
    balance -= 100;  
} else {  
    insufficient funds  
}
```

Thread #2

```
if (balance ≥ 50) {  
    balance -= 50;  
} else {  
    insufficient funds  
}
```

Memory = Shared Resource

balance := 100

Thread #1

```
if (balance ≥ 100) {  
    balance -= 100;  
} else {  
    insufficient funds  
}
```

a "critical section"
... at any moment, only
one of these... processes
is in its critical section

Dijkstra 1965

A black and white photograph showing a woman in a dark coat and hat working at a counter. Behind her, a long line of people, mostly women and children, wait patiently. The scene illustrates the concept of threads waiting for access to a shared resource.

Shared
Resource

Waiting
Threads

Thread in
Critical Region

a "critical section"

... at any moment, only
one of these... processes
is in its critical section

Dijkstra 1965

I will tell you how
to implement a lock.

But Dijkstra, how
do I implement this?

impressively, Dijkstra did this
without assuming any sort of
compare-and-swap instruction

We have locks, can we go home?

(or atomic test-and-set, or semaphores)

Where do I put locks? And what should happen if I leave out locks?

Thread #1

lock(1)

$x := x + 1$

$y := y + 1$

unlock(1)

Thread #2

$\text{tmp} := x + y$

Locks should be associated with the resource

How should I schedule requests?

Thread #1

```
lock(1)  
// high priority request  
unlock(1)
```

Thread #2

```
lock(1)  
//only run me if buffer  
// is empty  
unlock(1)
```

Schedule should be user-programmable

What about composability?

Thread #1

```
lock(1)  
lock(2)  
// critical region  
unlock(2)  
unlock(1)
```

Thread #2

```
lock(2)  
lock(1)  
// critical region  
unlock(1)  
unlock(2)
```

Must know which locks are
taken out and in
what order

Deadlock!

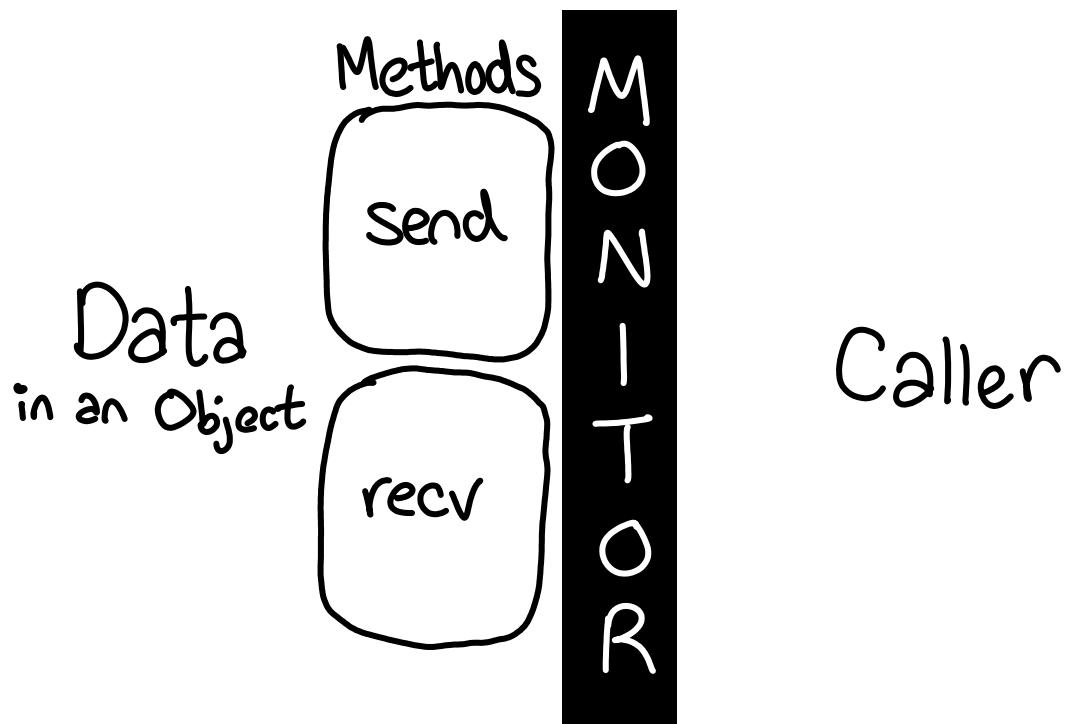
How can programming languages make concurrent programming easier?

By abstractions, of course!

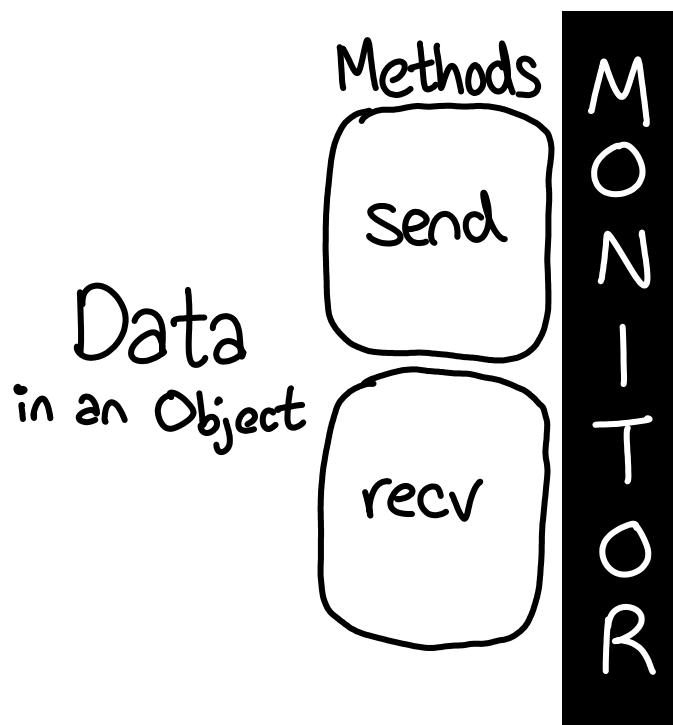
Monitors (Concurrent Pascal, Java)

```
public synchronized T  
recv() throws InterruptedException {  
    while (queue.isEmpty()) {  
        wait();  
    }  
    return queue.pop();  
}  
  
public synchronized void  
send(T x) { queue.push(x); notify(); }
```

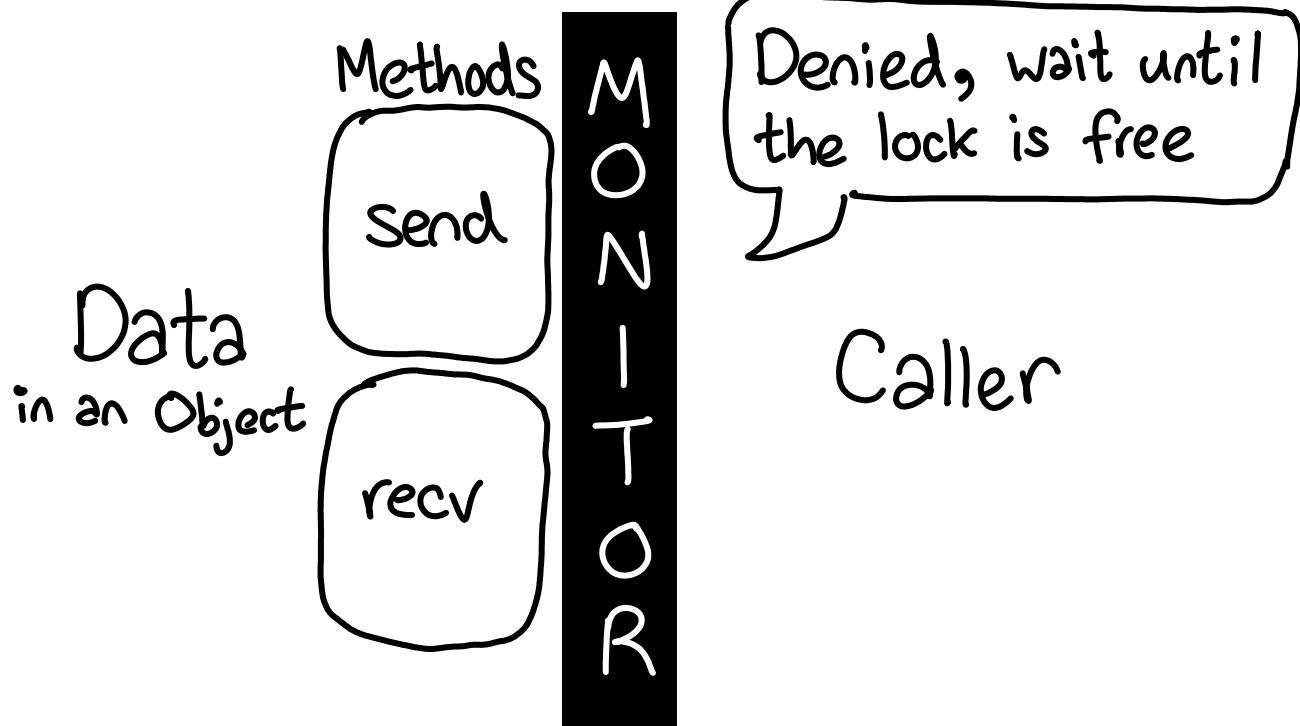
Monitors (Concurrent Pascal, Java)



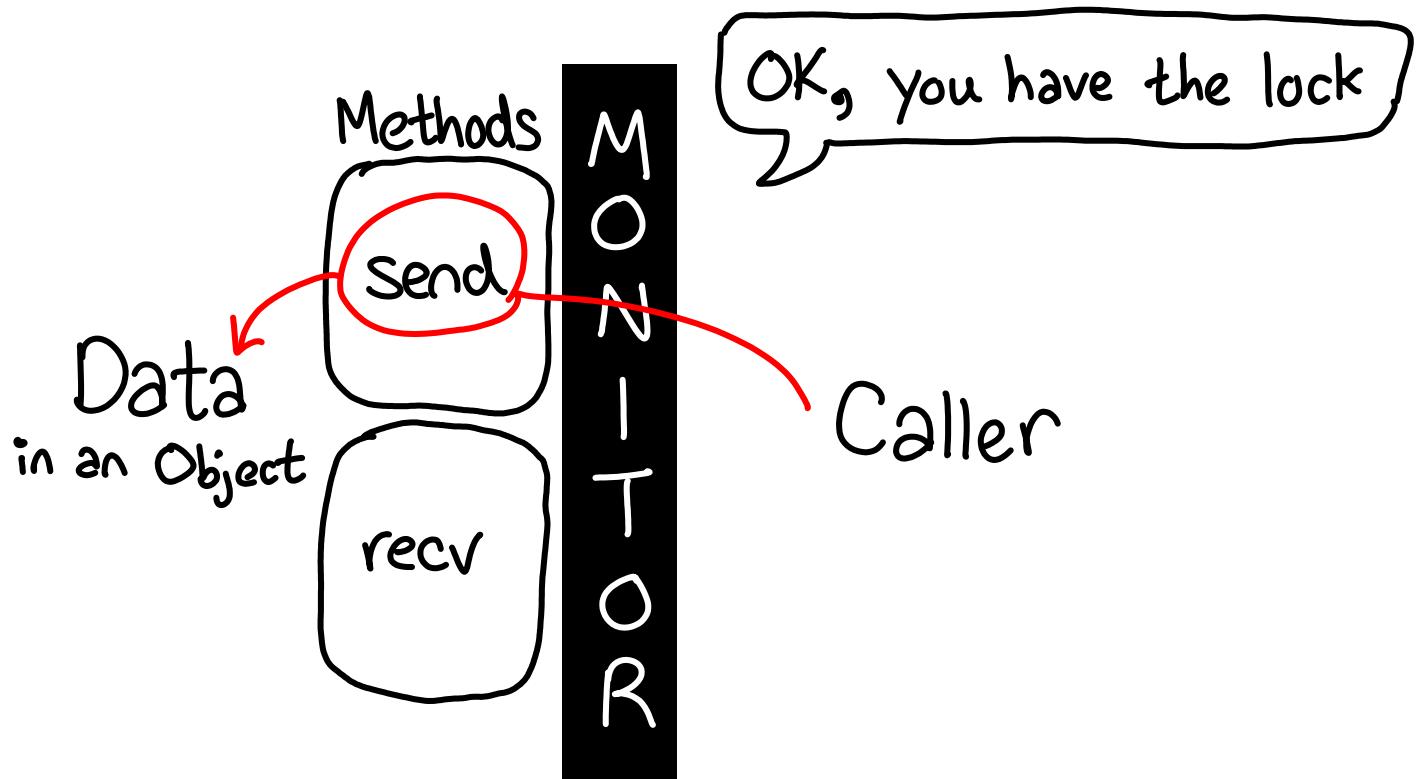
Monitors (Concurrent Pascal, Java)



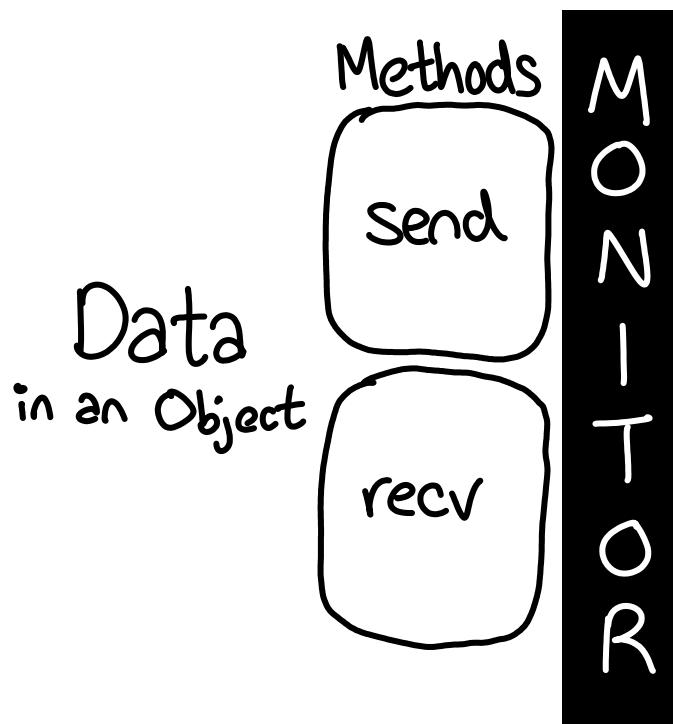
Monitors (Concurrent Pascal, Java)



Monitors (Concurrent Pascal, Java)



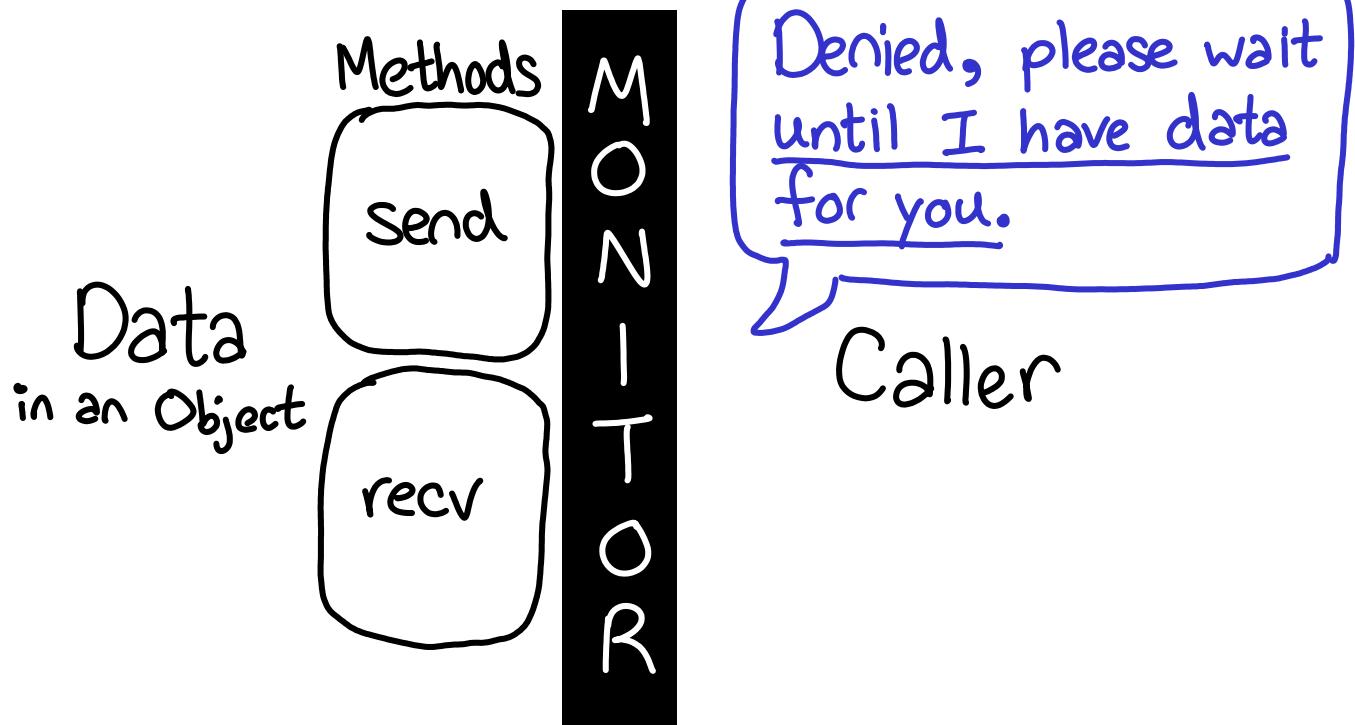
Monitors (Concurrent Pascal, Java)



I want to access
the data through
the recv method.

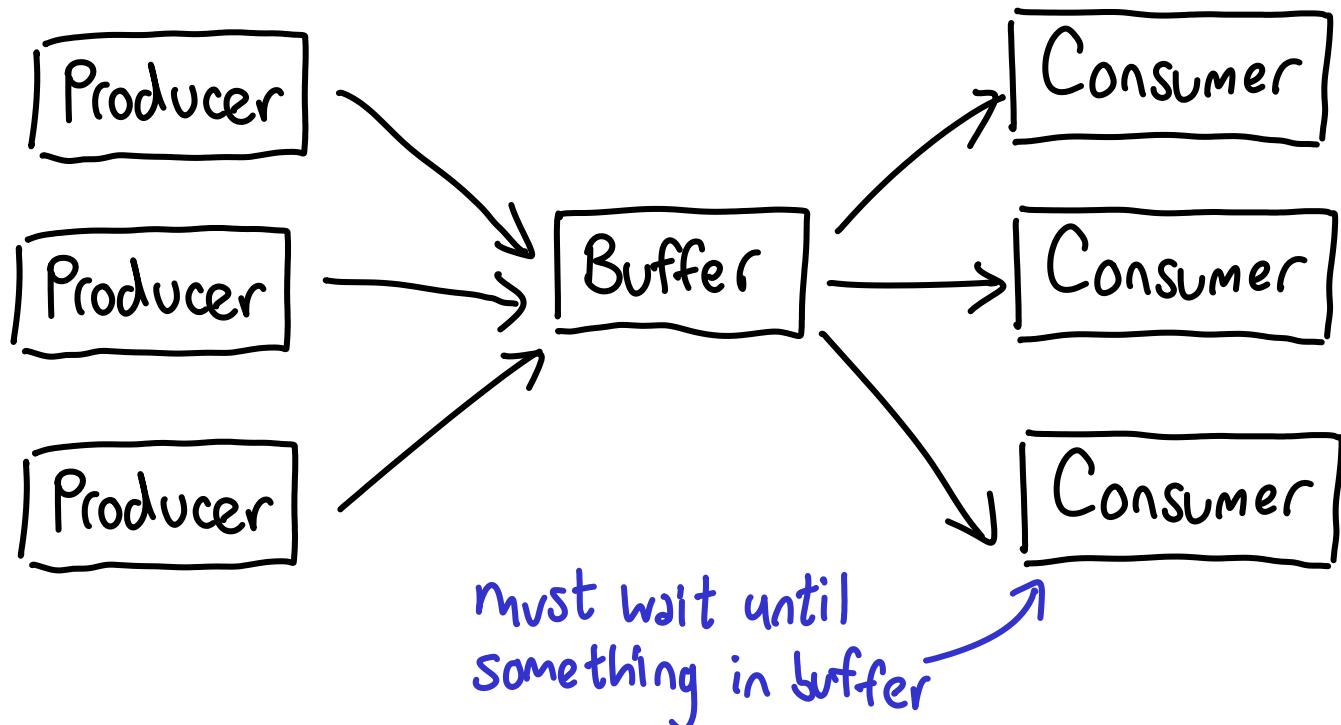
Caller

Monitors (Concurrent Pascal, Java)

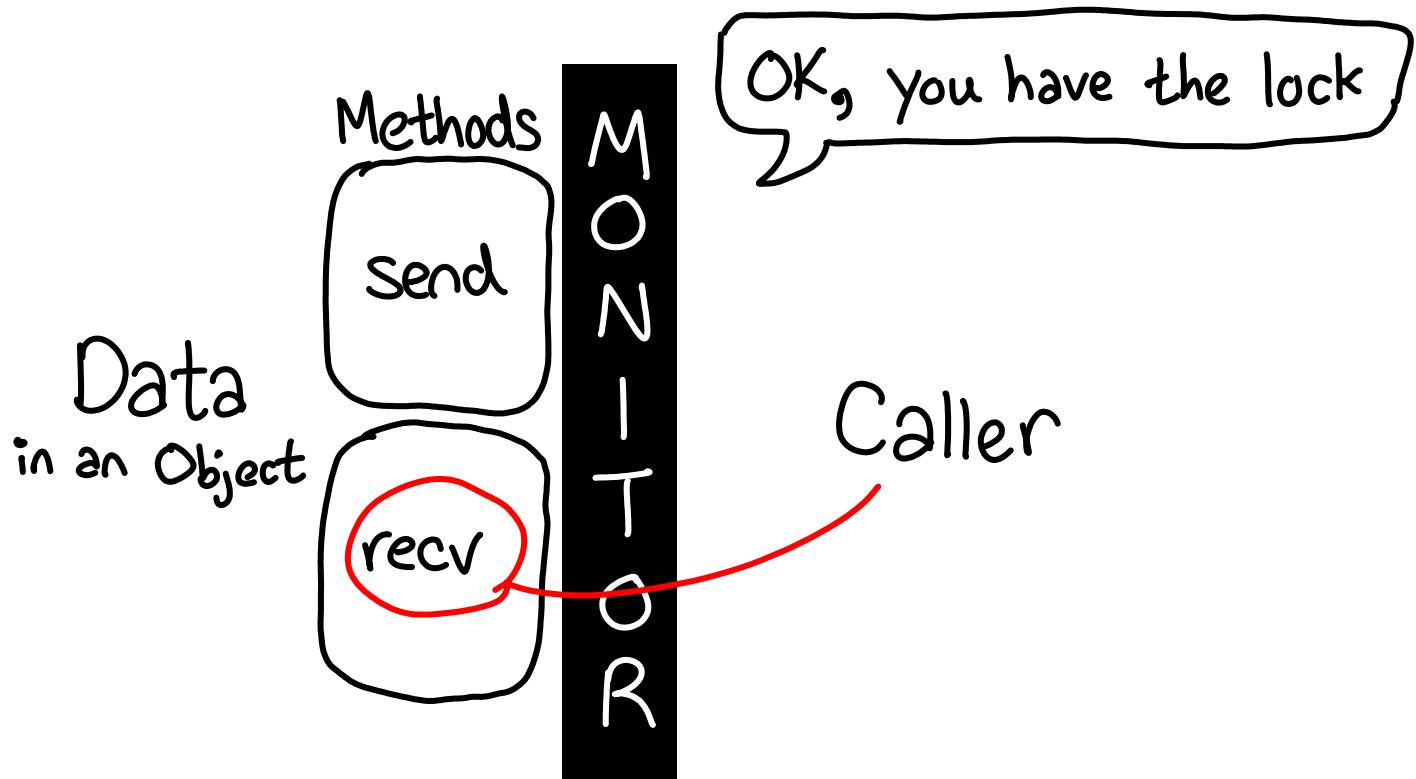


Monitors (Concurrent Pascal, Java)

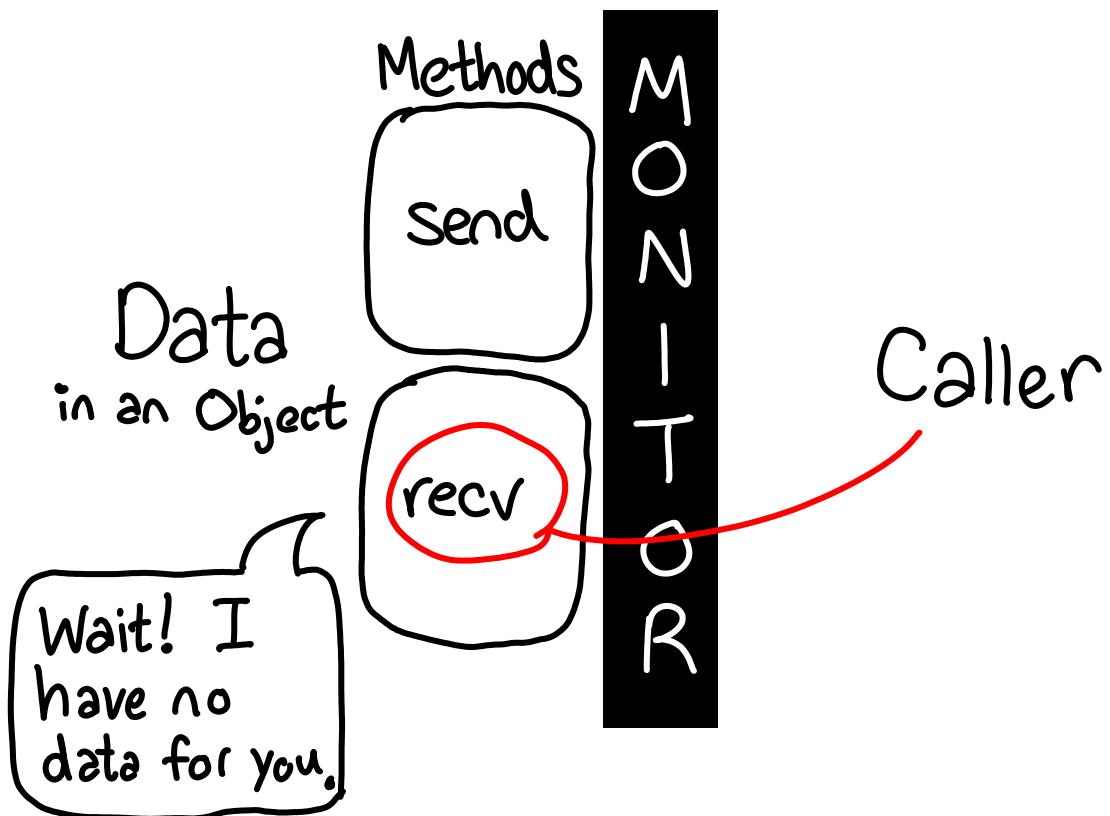
Producer-consumer



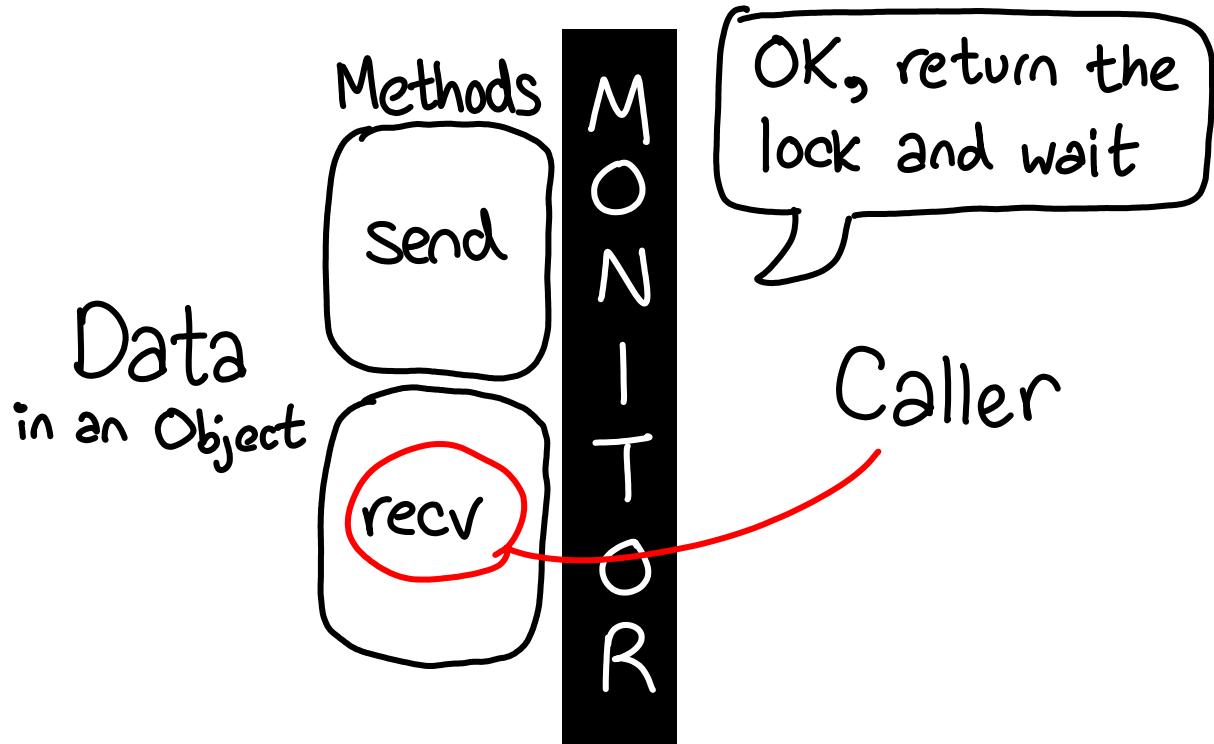
Monitors (Concurrent Pascal, Java)



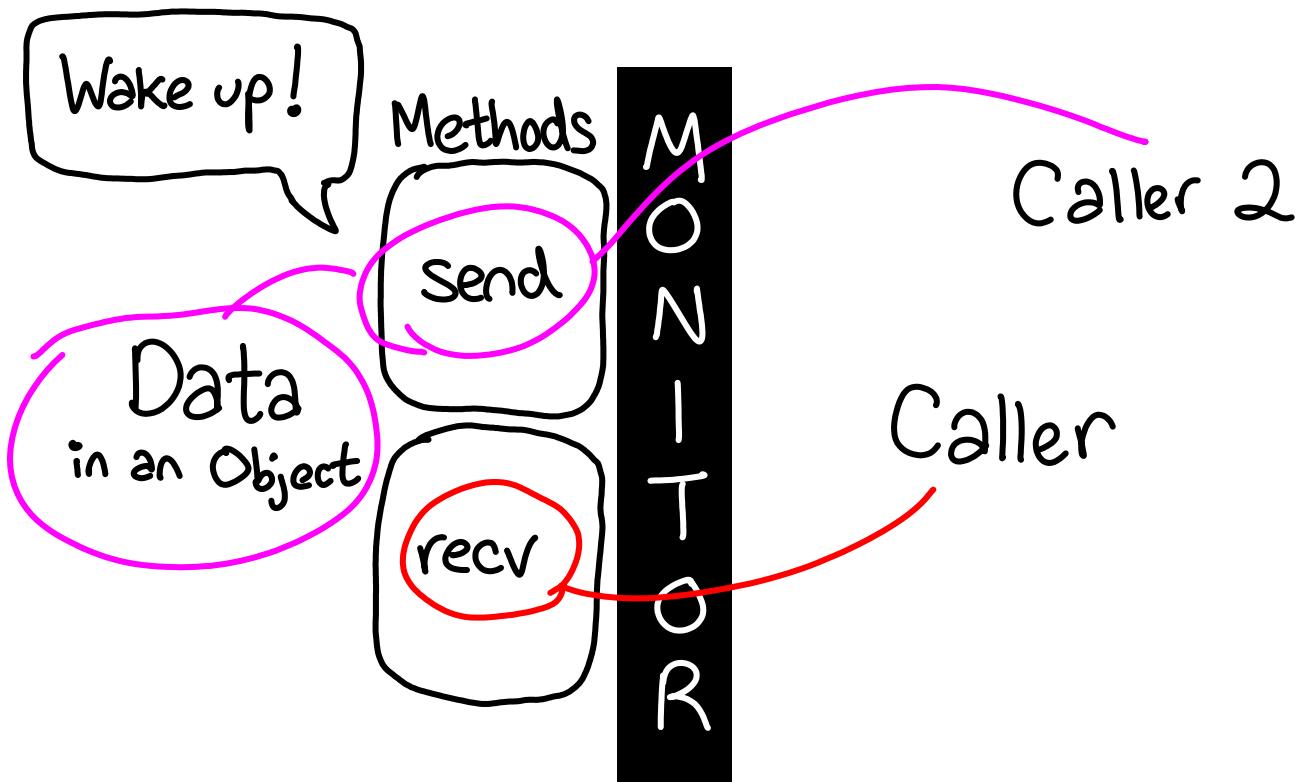
Monitors (Concurrent Pascal, Java)



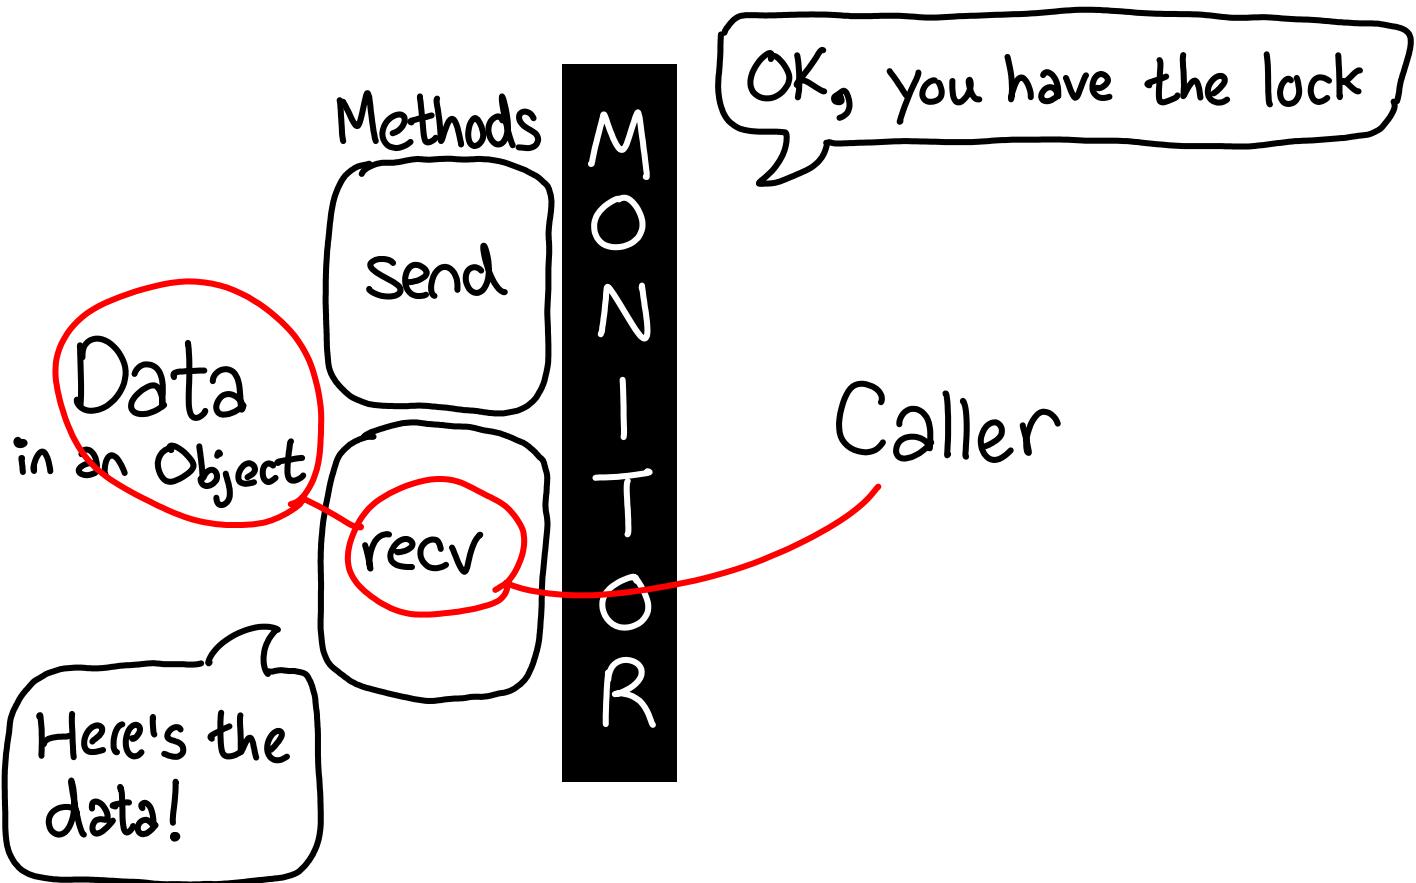
Monitors (Concurrent Pascal, Java)



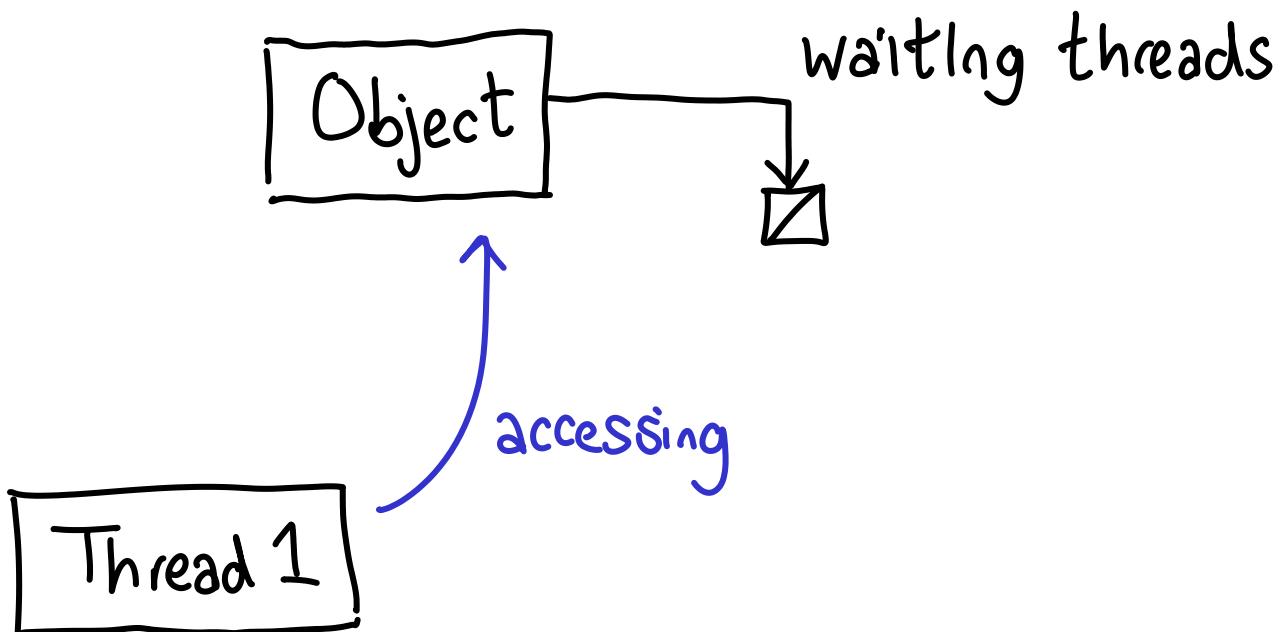
Monitors (Concurrent Pascal, Java)



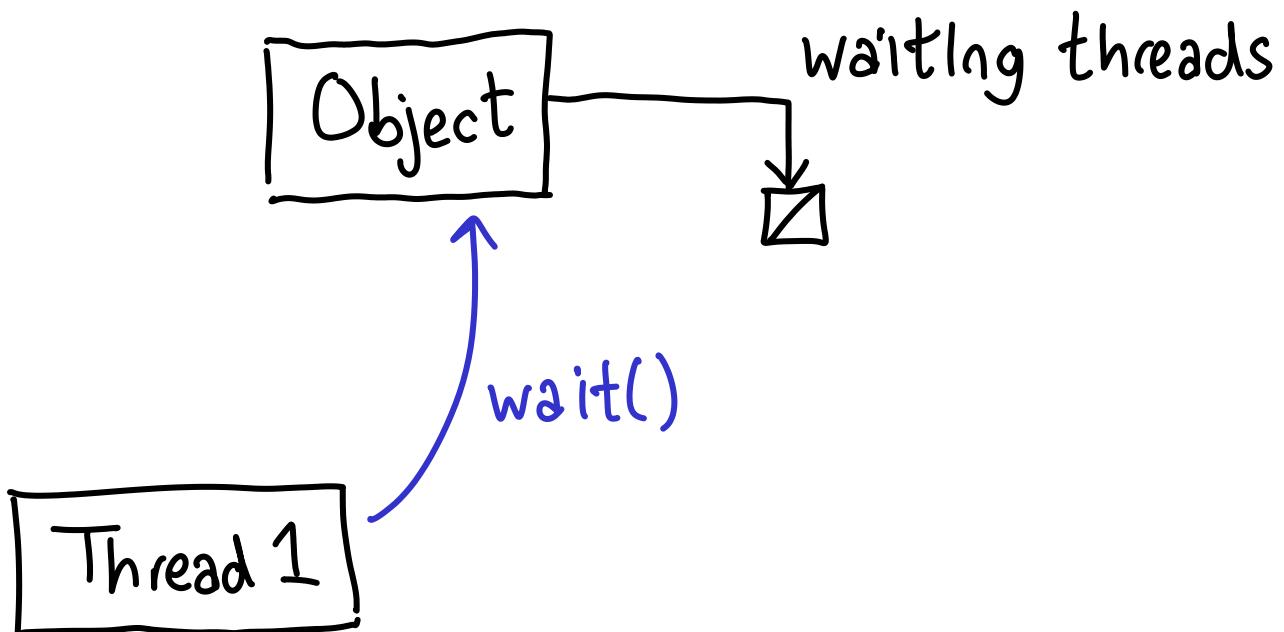
Monitors (Concurrent Pascal, Java)



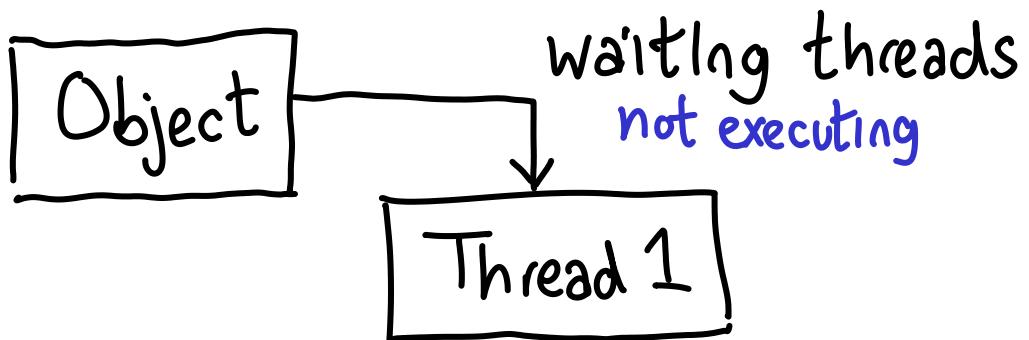
Monitors (Concurrent Pascal, Java)



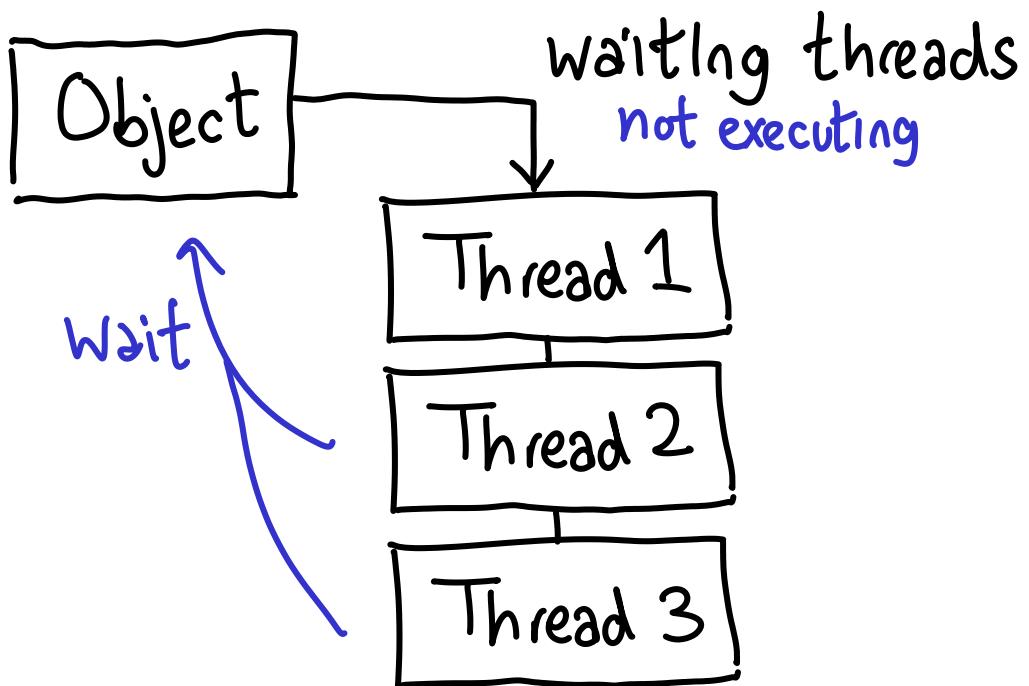
Monitors (Concurrent Pascal, Java)



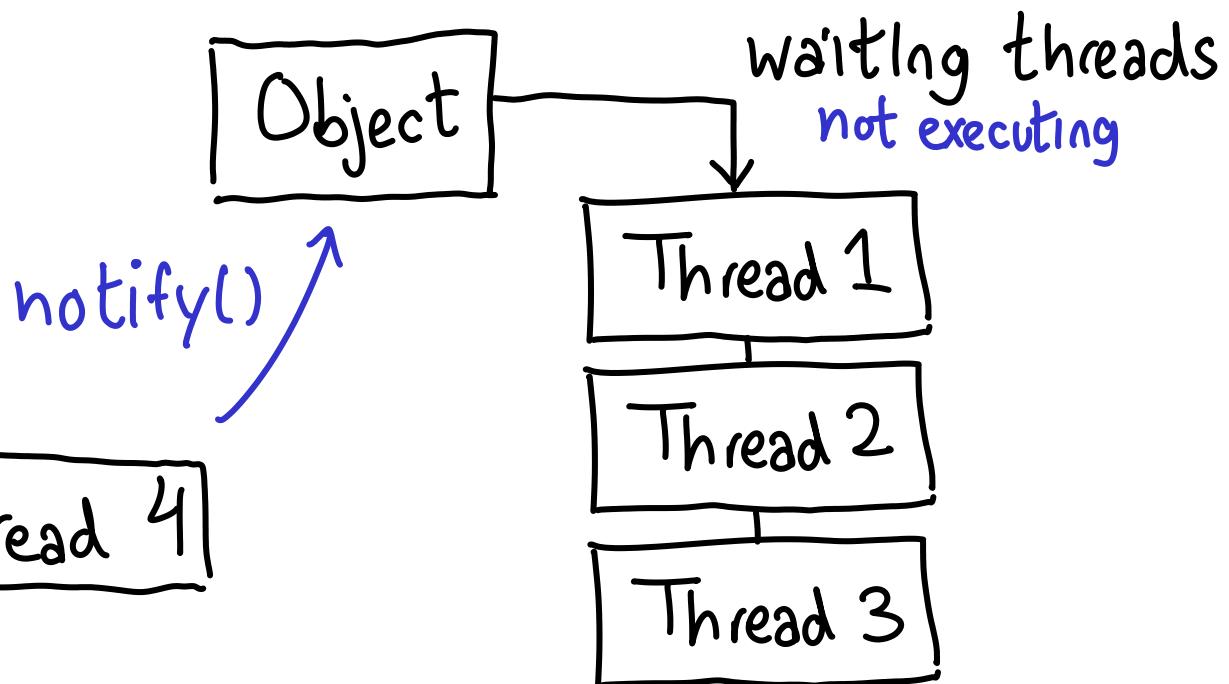
Monitors (Concurrent Pascal, Java)



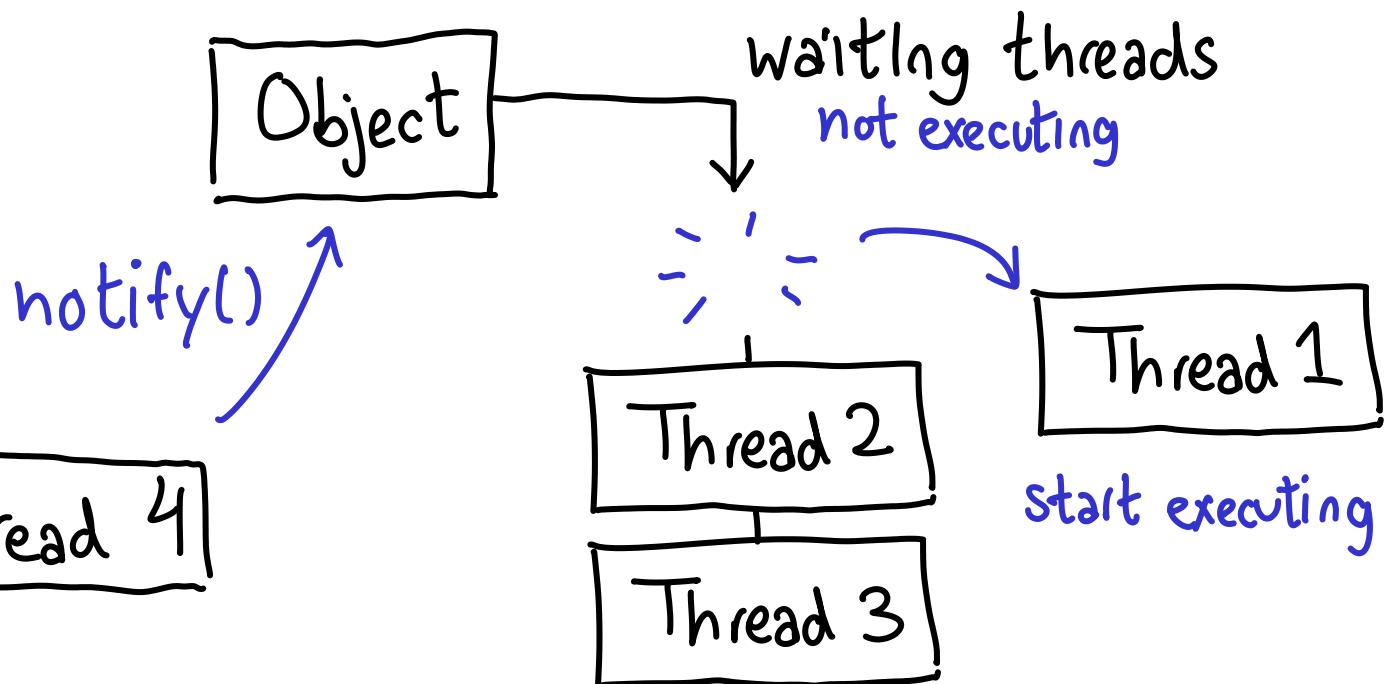
Monitors (Concurrent Pascal, Java)



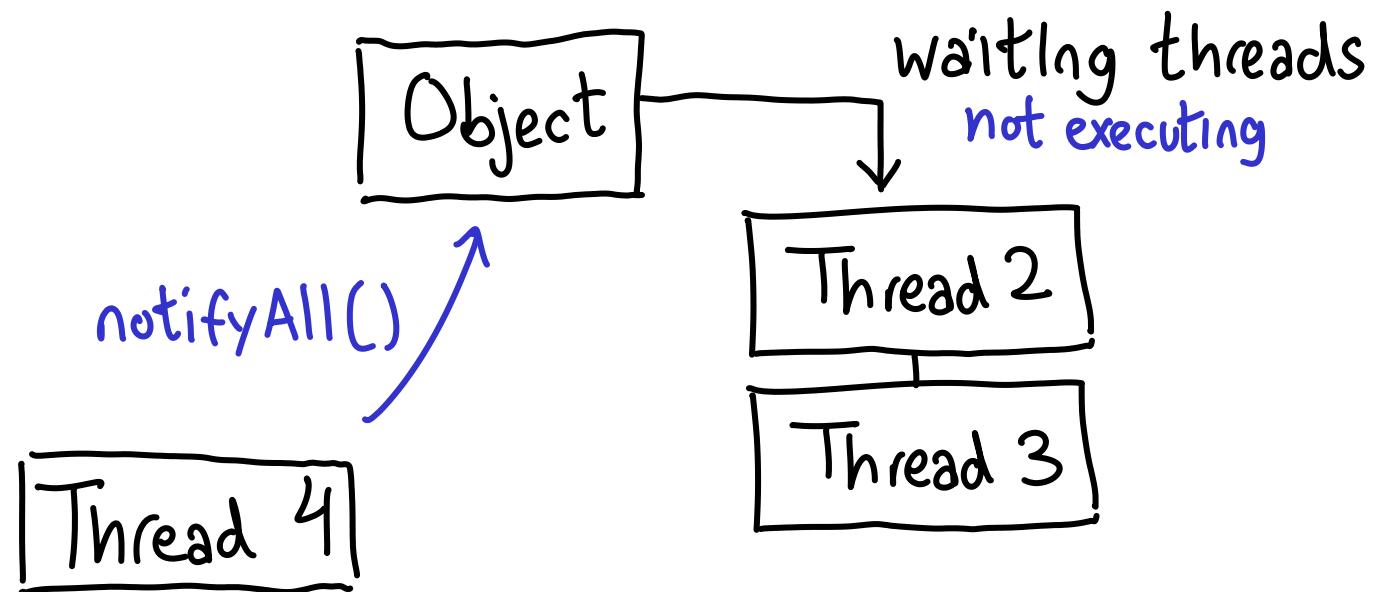
Monitors (Concurrent Pascal, Java)



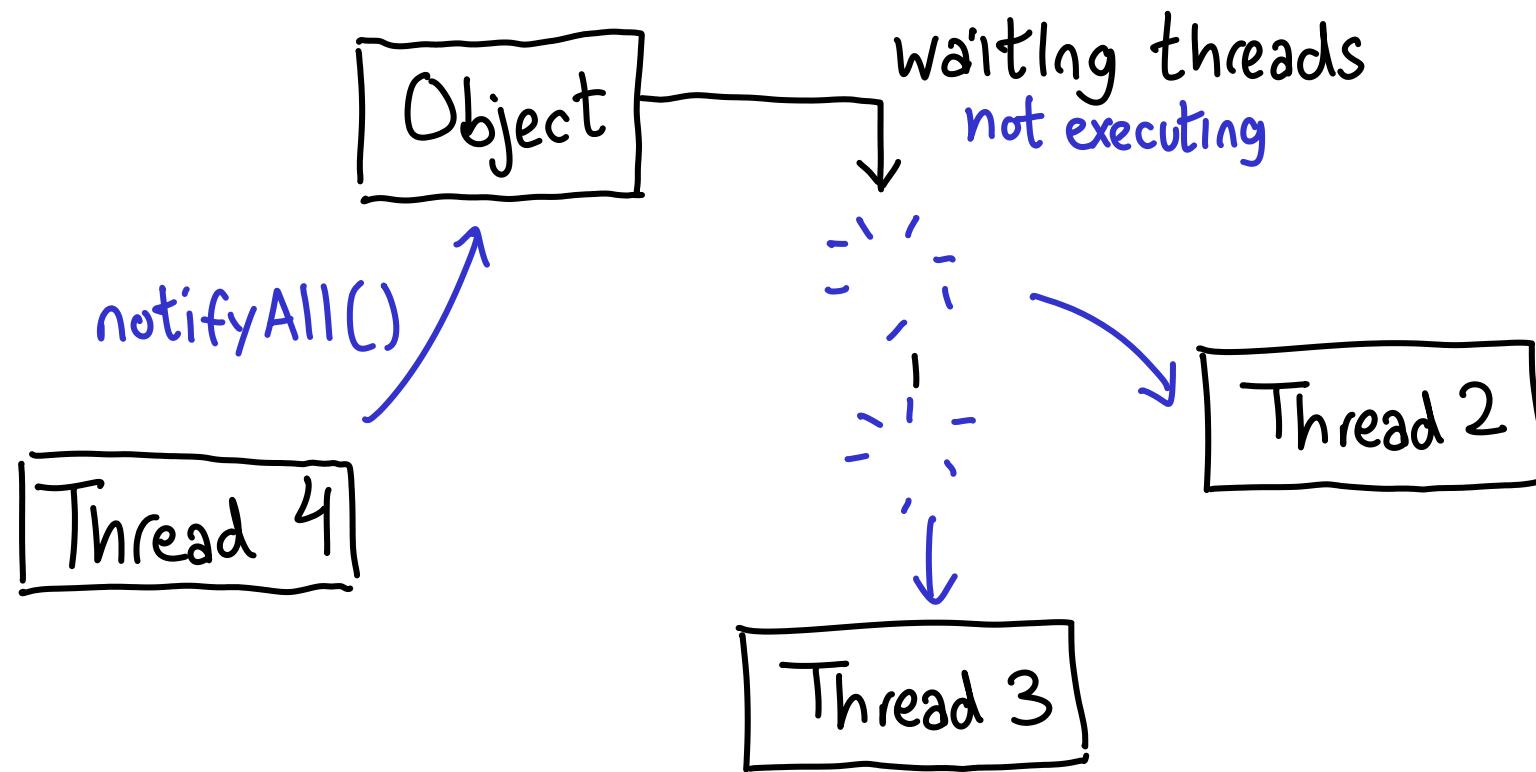
Monitors (Concurrent Pascal, Java)



Monitors (Concurrent Pascal, Java)



Monitors (Concurrent Pascal, Java)



Monitors (Concurrent Pascal, Java)

```
public synchronized T  
recv() throws InterruptedException {  
    while (queue.isEmpty()) {  
        wait();  
    }  
    return queue.pop();  
}  
  
public synchronized void  
send(T x) { queue.push(x); notify(); }
```

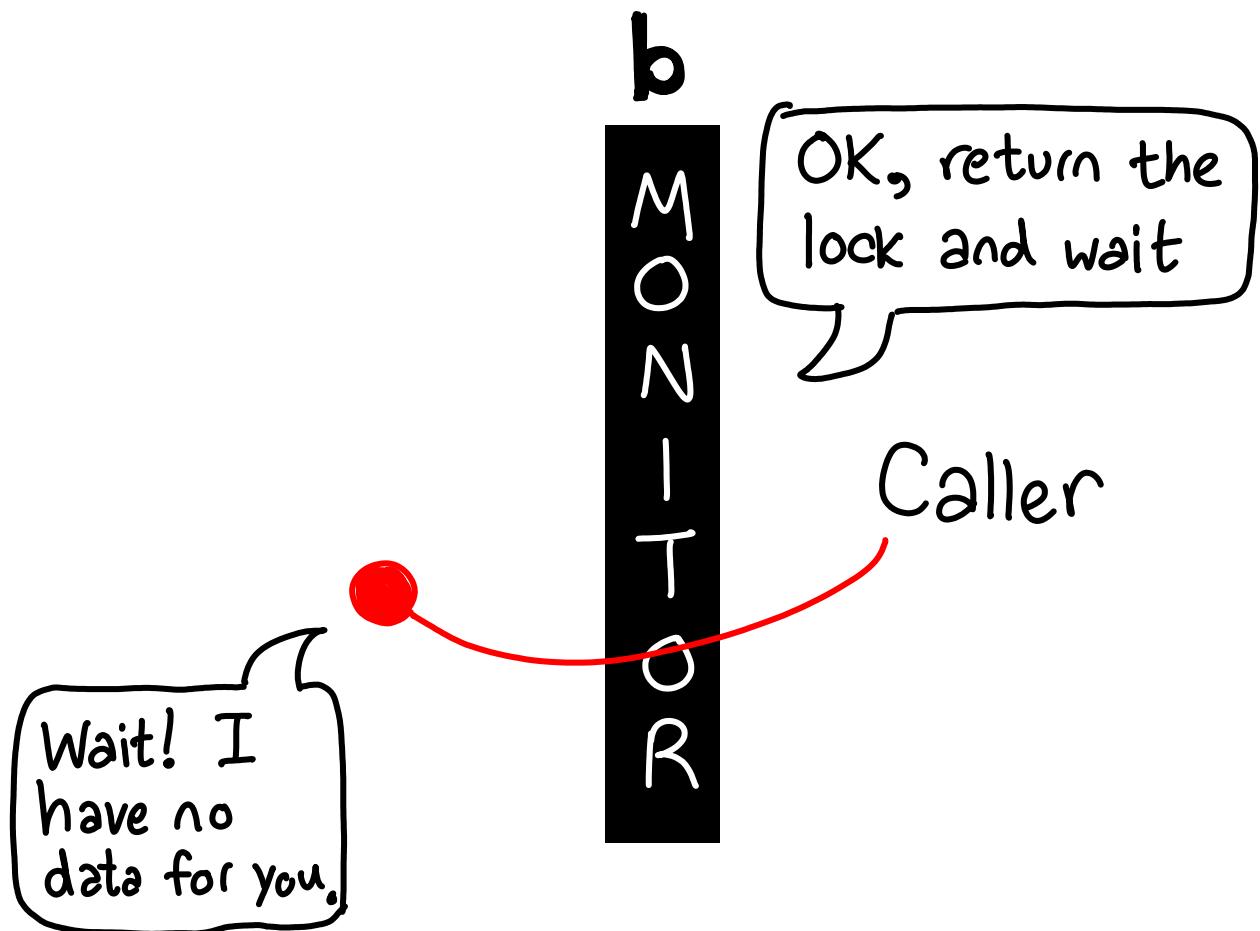
Monitors in Java

- Shared memory concurrency
programmer responsibility!
- Every object has a lock
- Synchronized / wait / notify / notifyAll
(Not part of type signature)

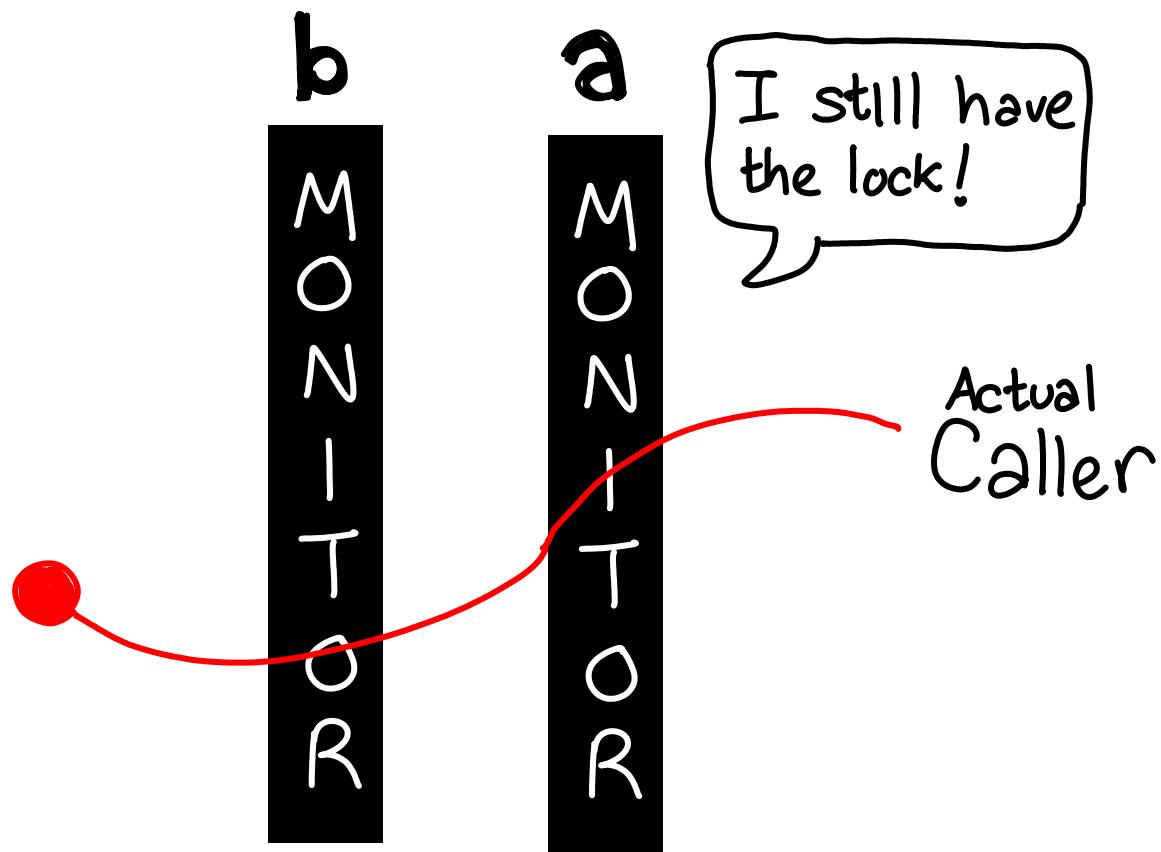
Monitors: Nested Monitor Lockout

```
Synchronized (a) {  
    Synchronized (b) {  
        a.wait();  
    }  
}
```

Monitors: Nested Monitor Lockout



Monitors: Nested Monitor Lockout

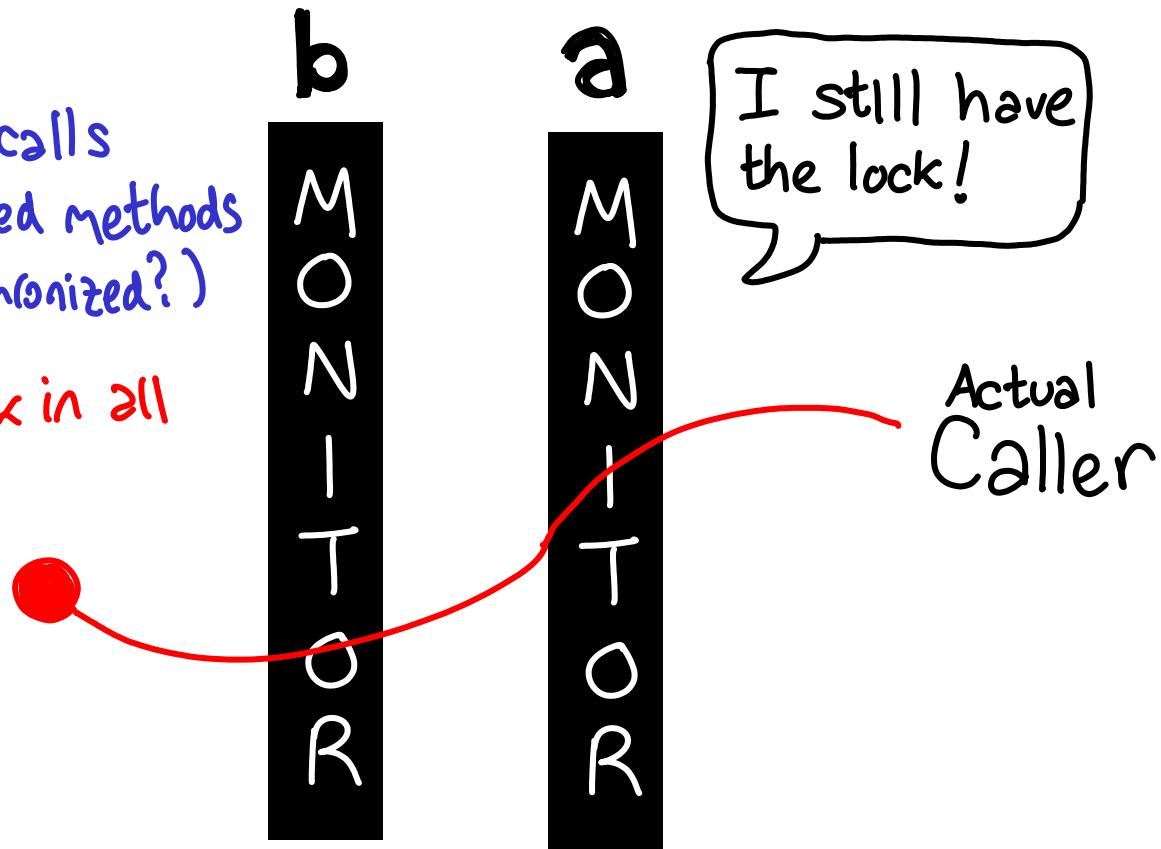


Monitors: Nested Monitor Lockout

"Advice":

No blocking calls
in synchronized methods
(use un-synchronized?)

Does not work in all
situations.



Monitors

synchronized / wait / notify

- ✓ Centralizes and hides synchronization logic associated with data
- ✗ A bit complicated / a bit inflexible
- ✗ Nested monitor lockout
- ✗ Lost notifications
- ✗ Priority must be supported (Java does not)
- ⚠ Doesn't solve deadlock

Message Passing

Monitor

I want to access
the data through
the recv method.

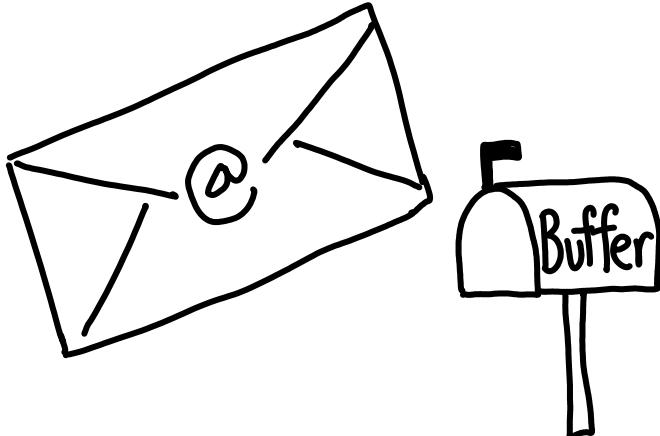
Caller

Denied, wait until
the lock is free

eliminate shared resources

Message Passing

Actor Model
Erlang



Asynchronous

Communicating
Sequential Processes
Go, Haskell

Send me the
data.

OK.

Synchronous

Message Passing: MVars in Haskell

one-place buffers

data MVar a

newEmptyMVar :: IO (MVar a)

newMVar :: a → IO (MVar a)

takeMVar :: MVar a → IO a

putMVar :: MVar a → a → IO a

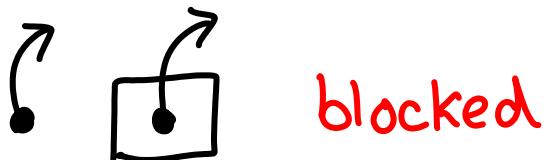
Message Passing: MVars in Haskell

one-place buffers

putMVar



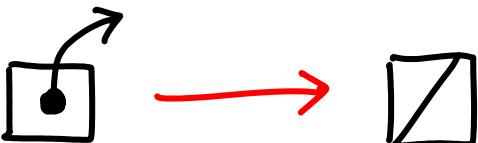
putMVar



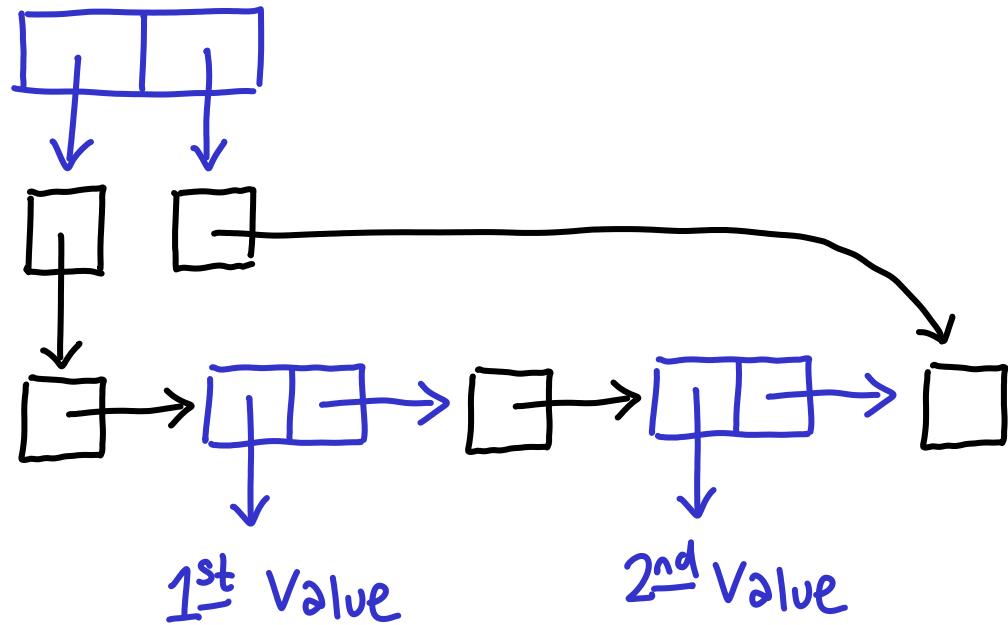
takeMVar



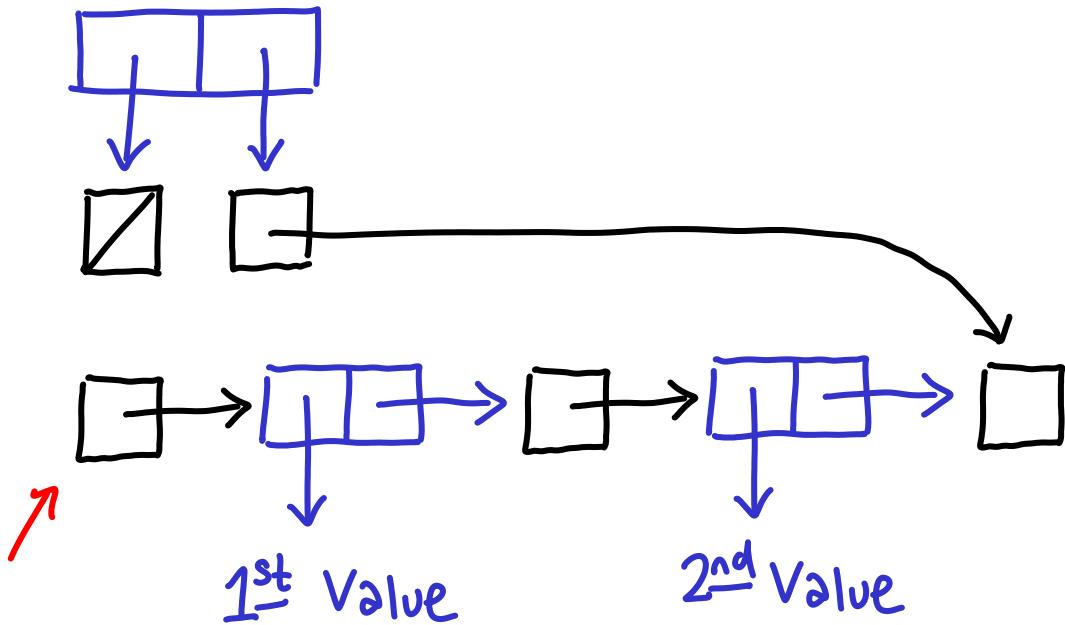
takeMVar



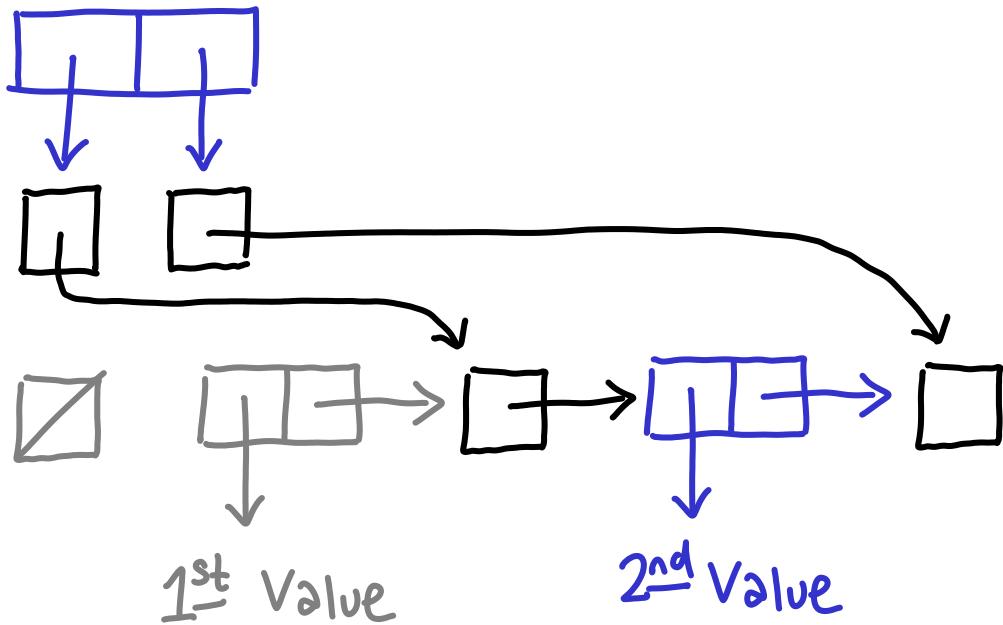
Message Passing: Channels with MVars



Message Passing: Channels with MVars



Message Passing: Channels with MVars



Message Passing

✓ Very simple (basis for many theoretical models;
also effective programming style)

for asynchronous messaging

- ? How big to maintain buffers?
- ? Unbounded nondeterminism

for synchronous messaging

- ⚠ Still easy to deadlock

Message passing and immutability

Once you send a message, must not r/w it!

Erlang separate address space

Haskell immutable data (but laziness!)

Rust/C++ ownership types (no aliasing)



like juggling chainsaws

Concurrency and immutability

Immutable data definitionally not shared

Deterministic Parallelism

Concurrency \neq Parallelism

Deterministic Parallelism

Concurrency

Parallelism

Availability

Speed

Interactivity

Distribution

Parallelism can be achieved with concurrency
Concurrency without parallelism

Deterministic Parallelism

Goal: Get speedups, w/o headaches of manual synchronization

- MapReduce (similar pipelines)
- Pure computation (sparks)
- Monotonic data structures (LVars)

Software Transactional Memory

(Wednesday)

Locks are **non-compositional**
(nested locking \rightsquigarrow deadlock)

```
synchronized (this) {  
    if (balance >= 100) {  
        balance -= 100;  
    } else {  
        insufficient funds  
    }  
}
```

locks

```
atomically {  
    if (balance >= 100) {  
        balance -= 100;  
    } else {  
        insufficient funds  
    }  
}
```

behave as if
it were an
atomic txn

Message Passing
Actors CSP

Transactional
Memory

Monitors
(Java)

Deterministic
Parallelism

Locks

Memory Model

Message Passing
Actors CSP

Transactional
Memory

Monitors
(Java)

Deterministic
Parallelism

Locks

data-races

Memory Model

Memory Model

Thread #1

$x = 1;$

$\text{tmp1} = y;$

Thread #2

$y = 1;$

$\text{tmp2} = x;$

What values of $\text{tmp1}/\text{tmp2}$ are allowed?

Memory Model

Thread #1

$x = 1;$

$\text{tmp1} = y;$

Thread #2

$y = 1;$

$\text{tmp2} = x;$

Java has 2 memory model

Thread #1

$\text{MOV } [x] \leftarrow 1$

$\text{MOV EAX} \leftarrow [y]$

Thread #2

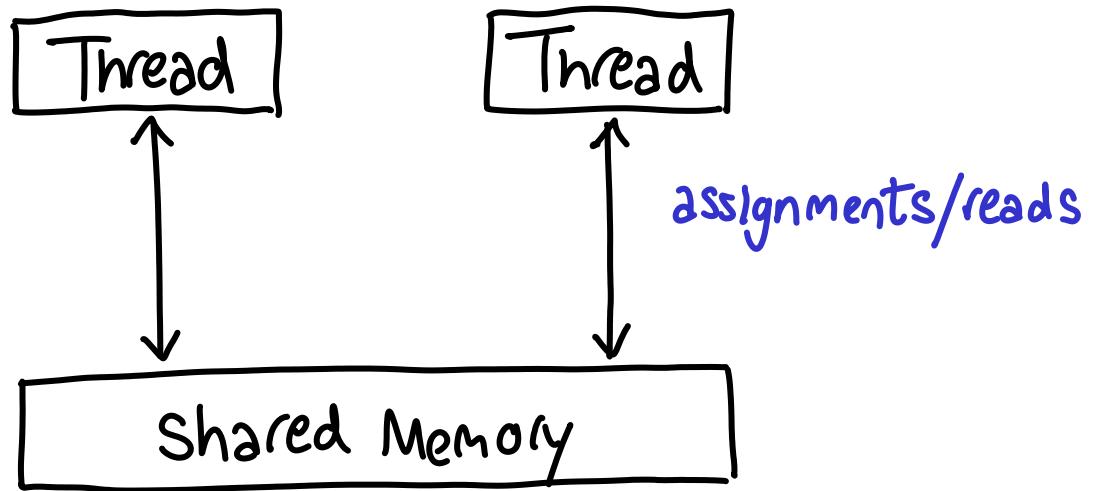
$\text{MOV } [y] \leftarrow 1$

$\text{MOV EBX} \leftarrow [x]$

so does your processor!

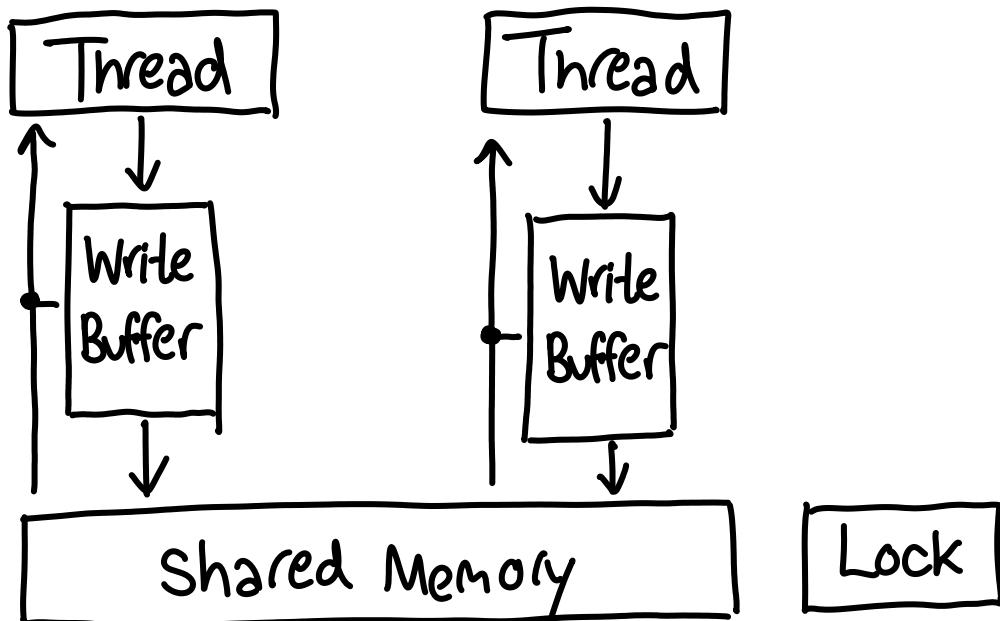
Memory Model: Sequential Consistency

Without data races, should look like this!



(implemented by basically no one)

Memory Model: Total Store Order



SPARC, x86

Memory Model: Total Store Order

Buffers are FIFO, serve reads if possible

MFENCE flushes buffer

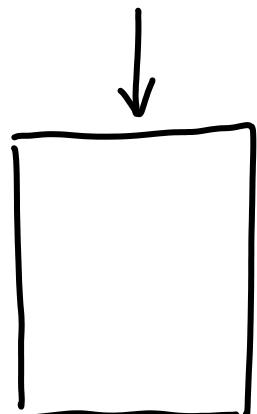
LOCK takes global lock, flushes buffer at end
(e.g. LOCKADD, XCHG)

Lock blocks buffered writes

Thread #1

MOV [x] ← 1

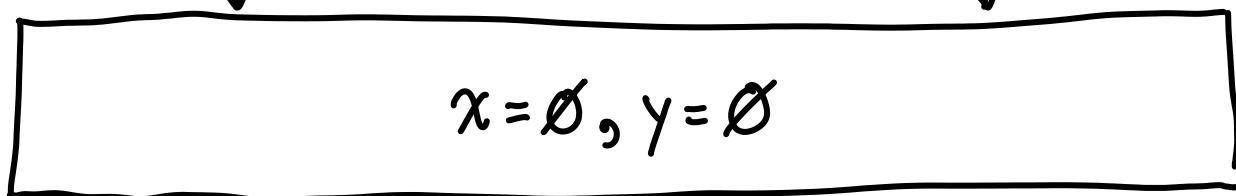
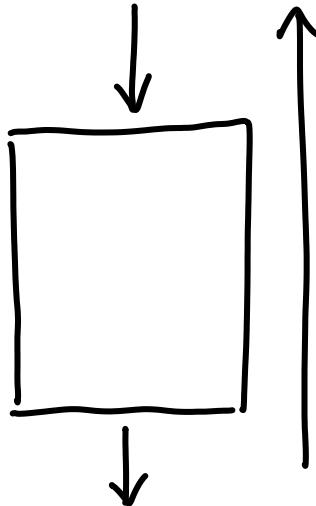
MOV EAX ← [y]



Thread #2

MOV [y] ← 1

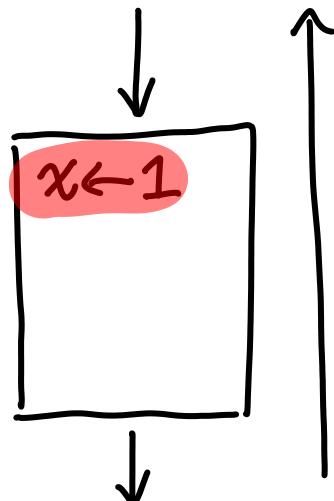
MOV EBX ← [x]



Thread #1

MOV [x] ← 1

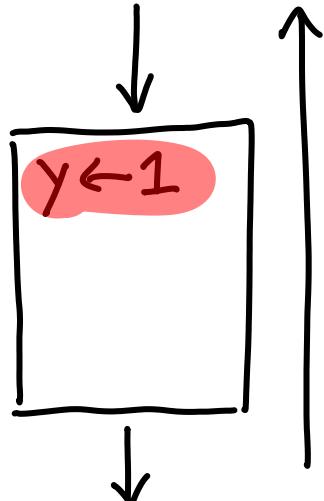
MOV EAX ← [y]



Thread #2

MOV [y] ← 1

MOV EBX ← [x]

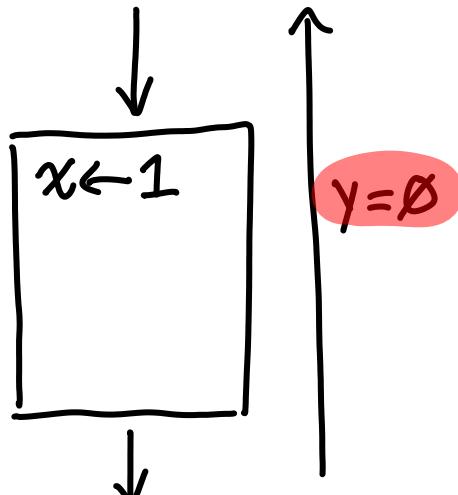


$x = \emptyset, y = \emptyset$

Thread #1

MOV [x] ← 1

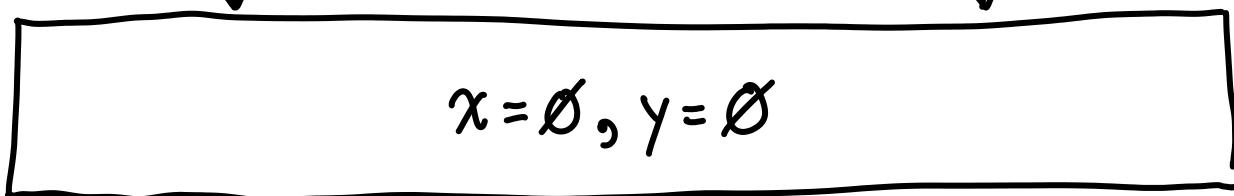
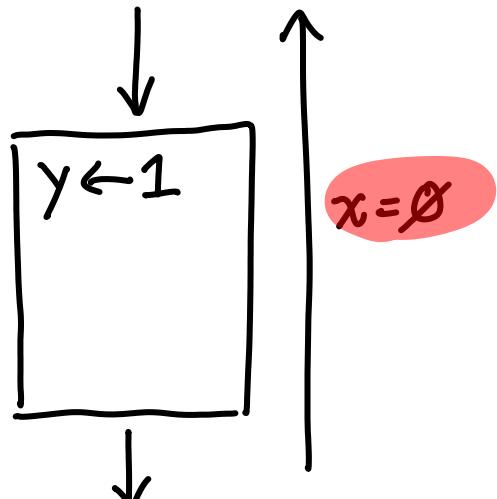
MOV EAX ← [y]



Thread #2

MOV [y] ← 1

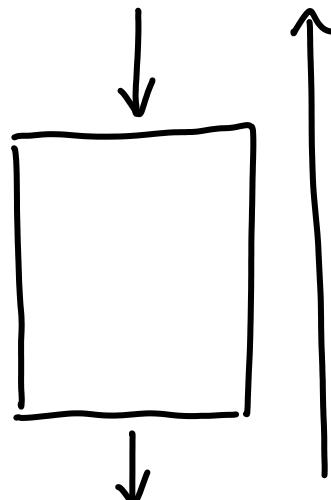
MOV EBX ← [x]



Thread #1

MOV [x] ← 1

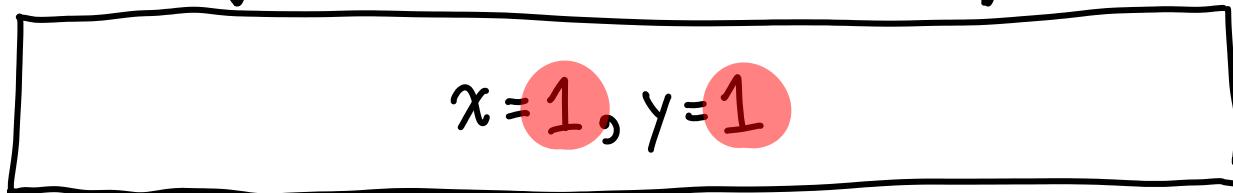
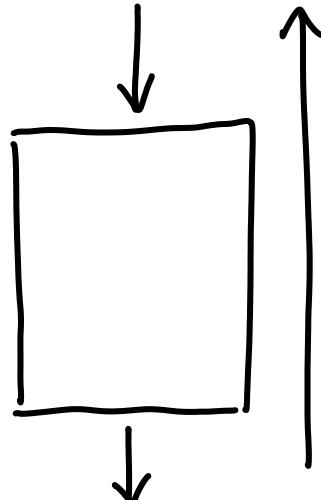
MOV EAX ← [y]



Thread #2

MOV [y] ← 1

MOV EBX ← [x]



On entry the address of spinlock is in register EAX
and the spinlock is unlocked iff its value is 1

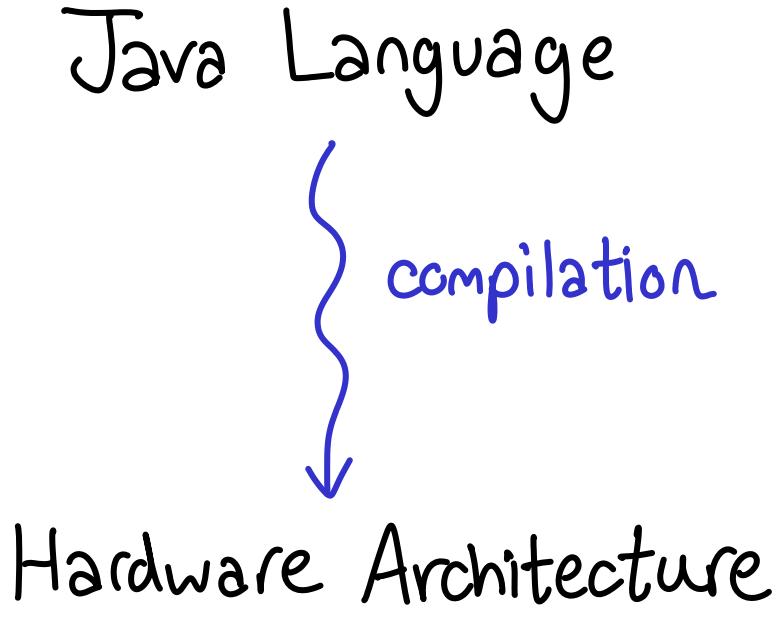
acquire:	LOCK;DEC [EAX]	; LOCK'd decrement of [EAX]
	JNS enter	; branch if [EAX] was ≥ 1
spin:	CMP [EAX],0	; test [EAX]
	JLE spin	; branch if [EAX] was ≤ 0
	JMP acquire	; try again
enter:	; the critical section starts here	
release:	MOV [EAX] $\leftarrow 1$	

 does this need to be locked?

Memory Model: ARM/POWER

- May read/write out of order
- Writes may not become visible to all threads at the same time point

Memory Model: Java

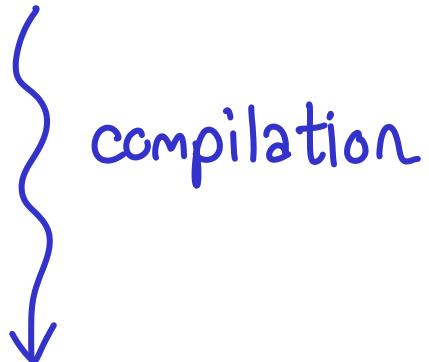


$x = 1;$
 $\text{tmp1} = y;$

$\text{MOV [x]} \leftarrow 1$
 $\text{MOV EAX} \leftarrow [y]$

Memory Model: Java

Java Language



Hardware Architecture

$x = 1;$
 $\text{tmp1} = y;$

optimizations!

MOV EAX $\leftarrow [y]$
MOV $[x] \leftarrow 1$

Memory Model: Java

- What optimizations are acceptable?

Notice, concurrency breaks abstraction barrier relied upon by optimization

- What program outputs can occur?

Undefined is unacceptable if it leads to a security problem

Memory Model: Java

Happens-before

$x=1;$ $y=2;$ *happens-before*

Memory Model: Java

Happens-before (Synchronizes-with)

Thread #1

LOCK M

X = 1

UNLOCK M

Thread #2

LOCK M

I = X

UNLOCK M

happens-before

Memory Model: Java

Happens-before (Synchronizes-with)

`volatile x;`

Thread #1

`x = 1;`

Thread #2

`y = x ;`

happens-before

(Not like C volatiles!)

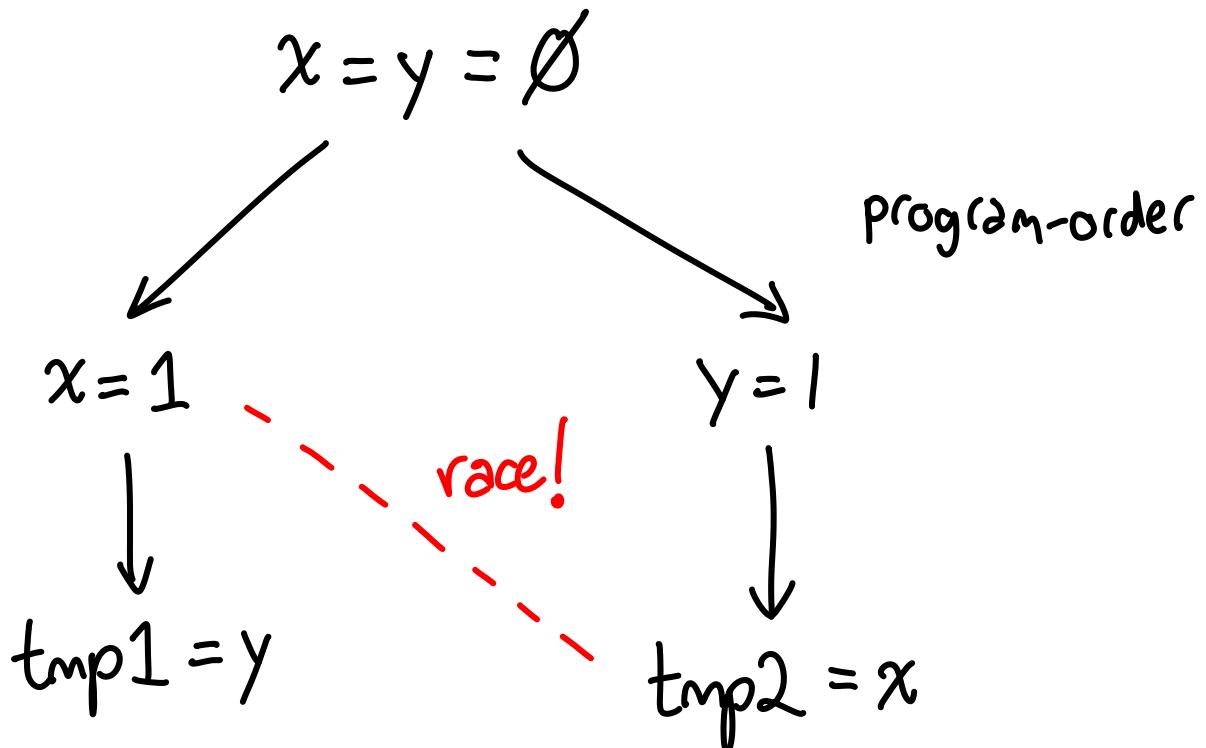
Memory Model: Java

Happens-before:

Program order + Synchronizes with

Data-race: two accesses on different threads
(one being a write) not ordered by
happens-before

If no data-races, program acts as if it is
sequentially consistent (if not? don't ask!)



Double-checked locking

```
private static Something instance = null;  
public Something getInstance() {  
    if (instance == null) {  
        synchronized (this) {  
            if (instance == null) {  
                instance = new Something();  
            }  
        }  
    }  
    return instance;  
}
```

Double-checked locking

READ instance

LOCK

READ instance

INITIALIZE Something

WRITE instance

UNLOCK

*Synchronizes
with*

data race

READ instance

LOCK

READ instance

INITIALIZE Something

WRITE instance

UNLOCK

thread 1

thread 2

Double-checked locking

READ instance

LOCK

READ instance

WRITE instance

INITALIZE Something

UNLOCK

reorder!

thread 2

reads the uninitialized object

READ instance

thread 1

Recall

How can programming languages make concurrent programming easier?

What abstractions are most effective?

Wednesday: Forward-looking research ideas

Message Passing
Actors CSP

Transactional
Memory

Monitors
(Java)

?

Say what we mean!

Deterministic
Parallelism

Locks

Memory Model

What does
it mean?

Extra examples

- Implementing monitors w/ MVars