

Atomicity and Software Transactional Memory

(thanks to SPJ for many slides)

The Problem


locks / condition variables

Traditional concurrency does

not compose well
deadlock / lost wakeups...

What we want


Layered
concurrency
abstractions




Concurrency primitives



What we have



Idea: Replace locks with atomic blocks



Atomic blocks
atomically/retry/orElse

Hardware

What's wrong with locks?

- Races forgot to lock
- Deadlock used wrong order
- Lost wakeups forgot to notify
- Error recovery forgot to cleanup on exception

Locks are non-compositional

```
class Account {  
    int balance;  
    synchronized void deposit(int amt) {  
        balance += amt; }  
    synchronized void withdraw(int amt) {  
        if (balance < amt)  
            throw new OutOfMoney();  
        balance -= amt; }  
}
```


add fund transfer?

Locks are non-compositional

```
class Account {  
    int balance;  
    synchronized void deposit(int amt) {  
        balance += amt; }  
    synchronized void withdraw(int amt) {  
        if (balance < amt)  
            throw new OutOfMoney();  
        balance -= amt; }  
    void transfer_wrong1(Account o, int amt) {  
        o.withdraw(amt);  
        this.deposit(amt);  
    }  
}
```

← race condition

Locks are non-compositional

```
class Account {  
    int balance;  
    synchronized void deposit(int amt) {  
        balance += amt; }  
    synchronized void withdraw(int amt) {  
        if (balance < amt)  
            throw new OutOfMoney();  
        balance -= amt; }  
     synchronized void transfer_wrong  
                                (Account o, int amt) {  
        o.withdraw(amt);  
        this.deposit(amt);  
    }  
}
```

deadlock

Limitations of race-freedom

```
class Ref {  
    int i;  
    void inc() {  
        int t;  
        synchronized (this) {  
            t = i;  
        }  
        synchronized (this) {  
            i = t + 1;  
        }  
    }  
}
```

Race free!

Doesn't do what
you want!

Race freedom does not
eliminate concurrency
errors.
(SC is still not easy!)

Limitations of race-freedom

```
class Ref {  
    int i;  
    void inc() {  
        int t;  
        synchronized (this) {  
            t = i;  
            i = t + 1;  
        }  
    }  
    int read() { return i; }  
}
```

Has a race!

Does what you want.

Race freedom is **not**
necessary to eliminate
errors!

Locks are absurdly hard to get right

Coding style

Difficulty of queue
implementation

Sequential code

Undergraduate

Locks are absurdly hard to get right

Coding style

Difficulty of queue
implementation

Sequential code

Undergraduate

Locks

Publishable result
at PODC¹

Locks are absurdly hard to get right

Coding style

Difficulty of queue implementation

Sequential code

Undergraduate

Locks

Publishable result
at PODC¹

Atomic blocks

Undergraduate

¹ Simple, fast, and practical non-blocking concurrent queue algorithms

Atomic: easier-to-use, harder to implement

```
void deposit(int x) {  
    synchronized(this) {  
        int tmp = balance;  
        tmp += x;  
        balance = tmp;  
    }  
}
```

semantics:

lock acquire/release

```
void deposit(int x) {  
    atomically {  
        int tmp = balance;  
        tmp += x;  
        balance = tmp;  
    }  
}
```

semantics: (behave as if)
no interleaved execution


Atomic memory transactions

atomically { ... sequential code ... }

- All or nothing! (error recovery)
- Isolated!
- No deadlocks! (no locks!)

How does it work?

atomically { ... <code> ... }



One possibility:

1. Execute <code> w/o any locks
2. Log each read/write in a thread-local transaction log
3. Writes go to log only, not memory

At the end, validate the log

- If valid, atomically commit changes to memory
- If invalid, re-run transaction from beginning

Software Transactional Memory

- Original paper/patent Tom Knight
- Many research/experimental implementations
C/C++, C#, Java, OCaml, Python, Scala
- Haskell STM!

Functional core simplifies STM concepts
(Transaction access of non-transactional memory?)

Some questions

- Java has locking & Condition variables

Some questions

- Java has ~~locking~~ & Condition variables
atomically

Some questions

- Java has ~~locking~~ & Condition variables
atomically blocking?

Some questions

- Java has ~~locking~~ & Condition variables
atomically blocking?
- Is programming with atomic really this easy?
efficiency?

STM in Haskell

Why Haskell?

Other languages:

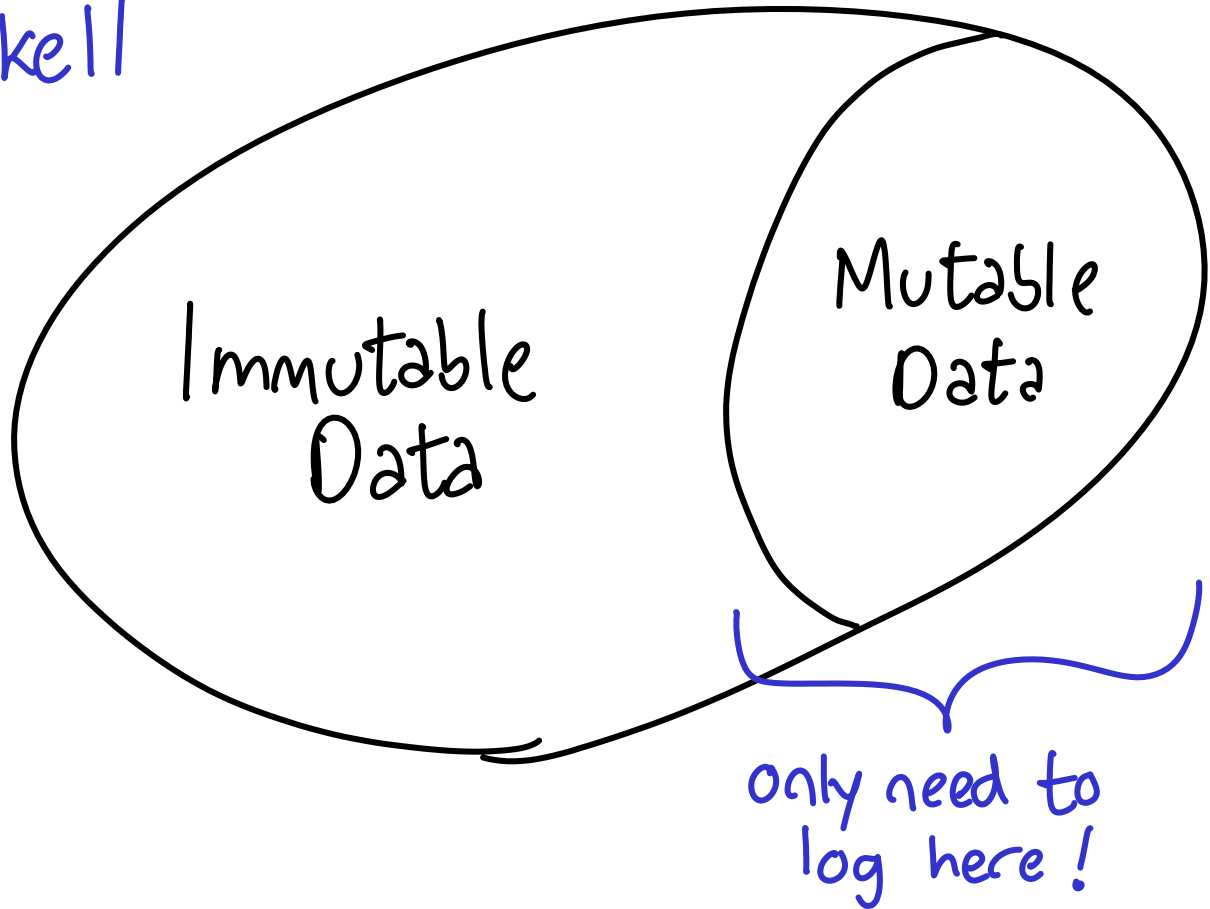


Mutable data

Logging = Expensive

Why Haskell?

Haskell





Why Haskell?

Transactions here only!

Pure code

IO code



Monads!

Immutable
Data

Mutable
Data

Why Haskell?



Recap: Effects in the type system

```
main = do { putStr (reverse "yes");  
            putStr "no" }
```

$:: \text{String}$ \swarrow no effects

$:: \text{IO } ()$ \nwarrow effects

Recap: Mutable State

reads & writes
100% explicit

```
main = do { r ← newIORef 0;  
           incr r;  
           s ← readIORef r;  
           print s }
```

$\text{incr} :: \text{IORef Int} \rightarrow \text{IO ()}$

```
incr r = do { v ← readIORef r;  
             writeIORef r (v+1) }
```

↑
r+1 disallowed!

Recap: Concurrency in Haskell

$\text{forkIO} :: \text{IO } () \rightarrow \text{IO ThreadId}$

```
main = do { r ← newIORef 0;  
           forkIO (incr r);  
           s ← readIORef r;  
           print s }
```

← race

(could fix with $\text{IORef} \rightarrow \text{MVar}$)

STM in Haskell

Idea: $\text{atomically} :: \text{IO } a \rightarrow \text{IO } a$ *almost!*

```
main = do { r ← newIORef 0;  
           forkIO (atomically (incr r));  
           s ← readIORef r;  
           print s }
```

Problem: shouldn't allow non-transactional access to r ..

STM in Haskell

↙ a monad!

Better:

- $\text{atomically} :: \text{STM } a \rightarrow \text{IO } a$
- $\text{newTVar} :: a \rightarrow \text{STM } (\text{TVar } a)$
- $\text{readTVar} :: \text{TVar } a \rightarrow \text{STM } a$
- $\text{writeTVar} :: \text{TVar } a \rightarrow a \rightarrow \text{STM}$

TVar can only be modified in a transaction!

can't modify
TVar outside of
transaction



```
main = do { r ← atomically (newTVar 0);  
           forkIO (atomically (incT r));  
           s ← atomically (readTVar r);  
           print s }
```

atomically is
just a function




$\text{incT} :: \text{TVar Int} \rightarrow \text{STM ()}$

```
incT r = do { v ← readTVar r;  
            writeTVar r (v+1) }
```



can't do
IO inside
transaction

Recap: Monads



$(\gg=) :: STM\ a \rightarrow (a \rightarrow STM\ b) \rightarrow STM\ b$

STM composes

Recap: Exceptions

$\text{throw STM} :: \text{Exception } e \Rightarrow e \rightarrow \text{STM } a$

$\text{catch STM} :: \text{Exception } e \Rightarrow \text{STM } a$

$\rightarrow (e \rightarrow \text{STM } a)$

$\rightarrow \text{STM } a$

If s throws exception in atomically s ,
abort transaction! No cleanup needed

Three new ideas:

retry orElse always

Idea 1: Compositional Blocking

$\text{withdraw} :: \text{TVar Int} \rightarrow \text{Int} \rightarrow \text{STM ()}$

$\text{withdraw acc } n =$

do { $\text{bal} \leftarrow \text{readTVar acc};$

if $\text{bal} < n$ then **retry**

else $\text{writeTVar acc (bal - n)}$ }

$\text{retry} :: \text{STM ()}$

Abort transaction and try again
from beginning. (Impl!)

Idea 1: Compositional Blocking

$\text{withdraw} :: \text{TVar Int} \rightarrow \text{Int} \rightarrow \text{STM ()}$

$\text{withdraw acc } n =$

do { $\text{bal} \leftarrow \text{readTVar acc};$

if $\text{bal} < n$ then **retry**

else $\text{writeTVar acc (bal - n)}$ }

- No condition variables
- Retrying thread woken up automatically
- No danger of forgetting to re-test conditions

Why is retry compositional?

Retry can occur anywhere in an atomic block

atomically (withdraw a1 3 >>
 withdraw a2 7)

← waits on
both
conditions

Non-composition alternative:

declare all conditions upfront

Idea 2: Choice

Suppose we want to transfer three dollars from either account a_1 or a_2 to account b

to account b

```
atomic ((withdraw a1 3  
        'orElse'  
        withdraw a2 3) >>  
        deposit b 3)
```

try this

...and if it retries

... then do this.

$$\text{orElse} :: \text{STM } a \rightarrow \text{STM } a \rightarrow \text{STM } a$$

Choice composes too!

transfer2 a1 a2 b

'orElse'

transfer2 a3 a4 b

It's associative! (but not commutative)

An algebra!

$$\text{retry 'orElse'} s \equiv s$$

$$s \text{ 'orElse' retry} \equiv s$$

MonadPlus!

Idea 3: Invariants

Goal: Establish invariants which are true on entry & exit from atomic

always :: STM Bool \rightarrow STM ()

newAccount :: STM (TVar Int)

newAccount = do

 v \leftarrow newTVar 0

 always (do cts \leftarrow readTVar v
 return (cts \geq 0))

 return v

arbitrary
STM code



Idea 3: Invariants

Goal: Establish invariants which are true on entry & exit from atomic

always:: STM Bool \rightarrow STM ()

- Adds a new invariant to pool of invariants
- Conceptually: checked after every txn
(Actually, only check for modified TVars;
garbage collect based on dead TVars)

spec?

See Composable Memory Transactions
for details!

GHC ships w/ a complete STM impl
import Control.Concurrent.STM

Microbenchmarks:

~50-80 ns

~20 ns

TVar read/put

MVar take/put

(mutex cost)

Worse: readTVar $O(n)$ in number of TVars in txn

STM in mainstream languages?

```
class Account {  
    int balance;  
    atomic void deposit(int amt) {  
        balance += amt; }  
    atomic void withdraw(int amt) {  
        if (balance < amt)  
            throw new OutOfMoney();  
        balance -= amt; }  
    atomic void transfer (Account o, int amt) {  
        o.withdraw(amt);  
        this.deposit(amt);  
    }  
}
```

STM in mainstream languages?

Trouble: type system doesn't control effects

Weak atomicity

Non-transactional code may see inconsistent states on transactional code.

Strong atomicity

Non-transactional code guaranteed to see a consistent state (performance hit!)

Enforcing Isolation and Ordering in STM

Performance?

- Naïve STM is hopelessly inefficient (think 6x slowdown or more)
- HTM: hardware transactional memory; hardware support for "lock elision"
- Hybrid STM: use HTM support to implement STM efficiently

↑ research direction

Atomicity not a silver bullet

How big should atomic blocks be? Races vs. Starvation

Thread 1

[atomic { x = 1; }
atomic { if (y == 0) retry; }


Thread 2

atomic {
if (x == 0) retry;
y = 1;
}

↑
incorrect to make this a
single atomic block

Atomicity not a silver bullet

STM can't deadlock, but it's not necessarily fair




Some parting thoughts

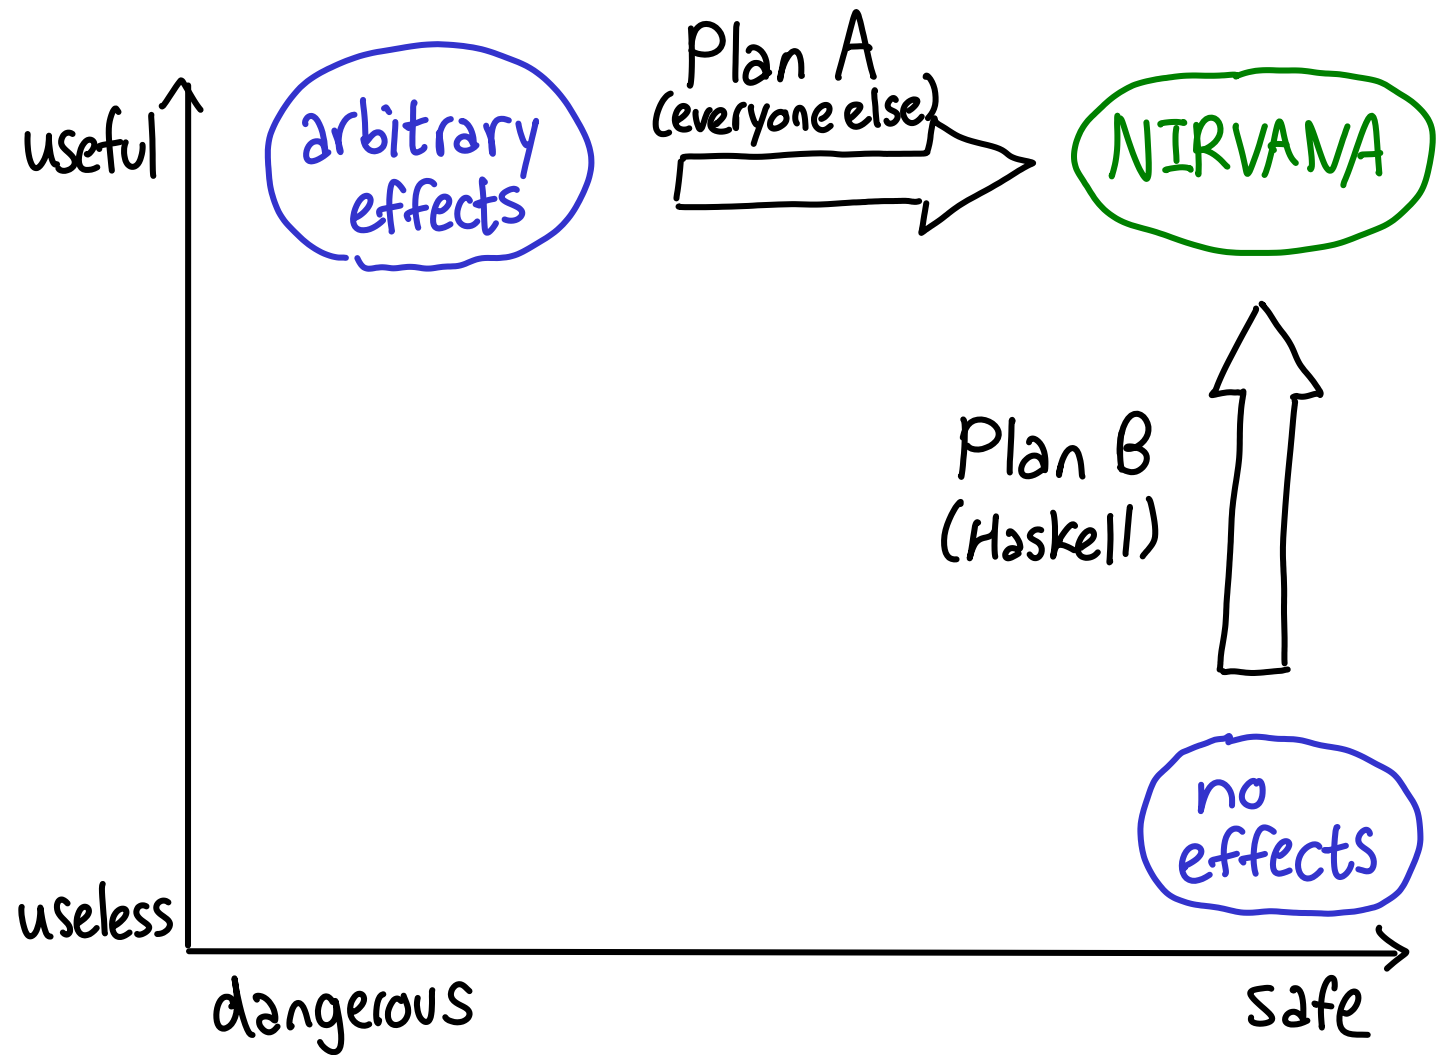
Side-effectful IO computation



```
graph TD; A([Side-effectful IO computation]) --- B([Pure effect free computation]); B --> A;
```

Pure effect free
computation







Default: any effect
Plan: add restrictions

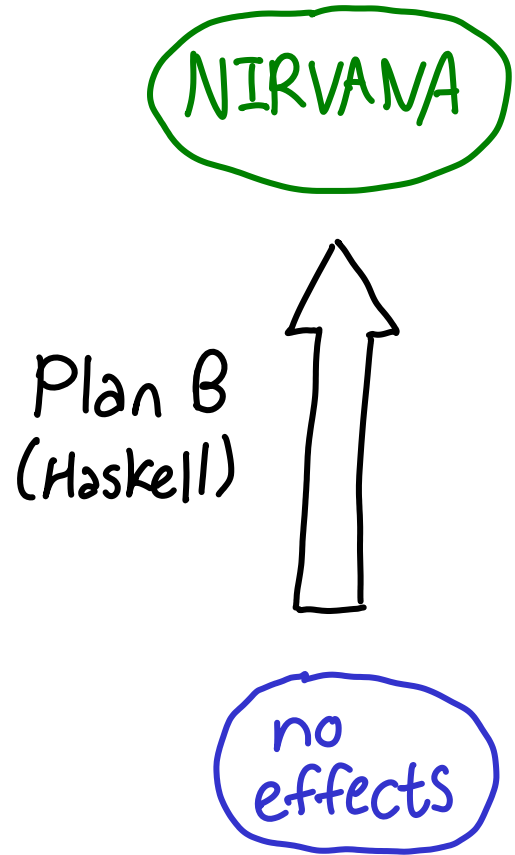
Examples: Regions
Ownership types
Vault, Spect#, Cyclone

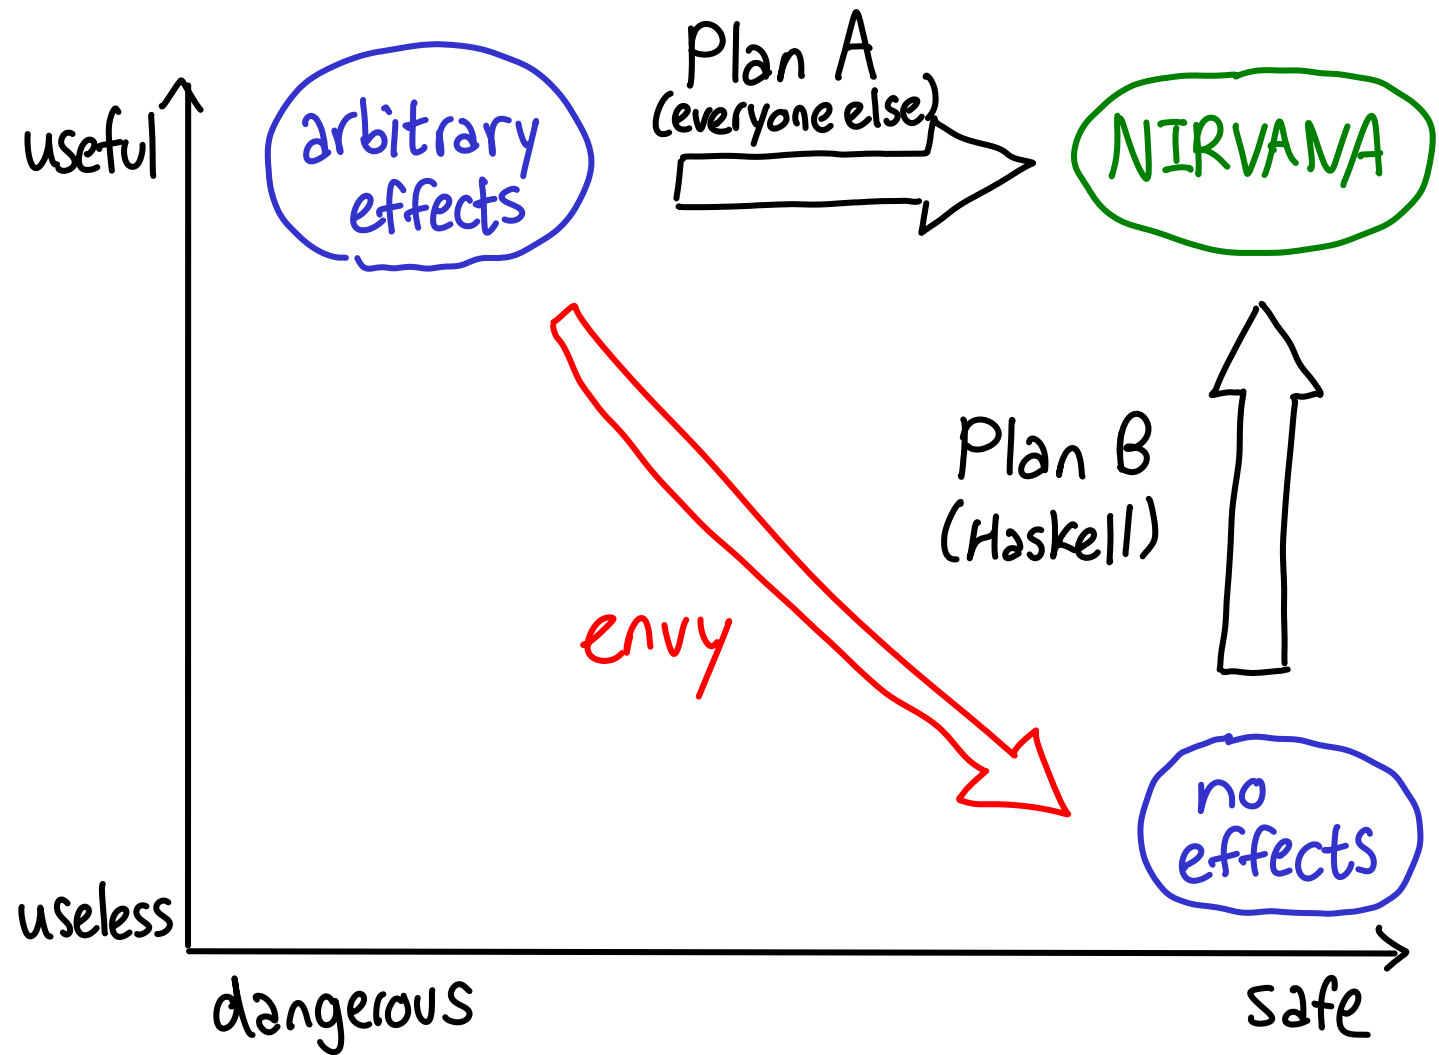
Default: no effects

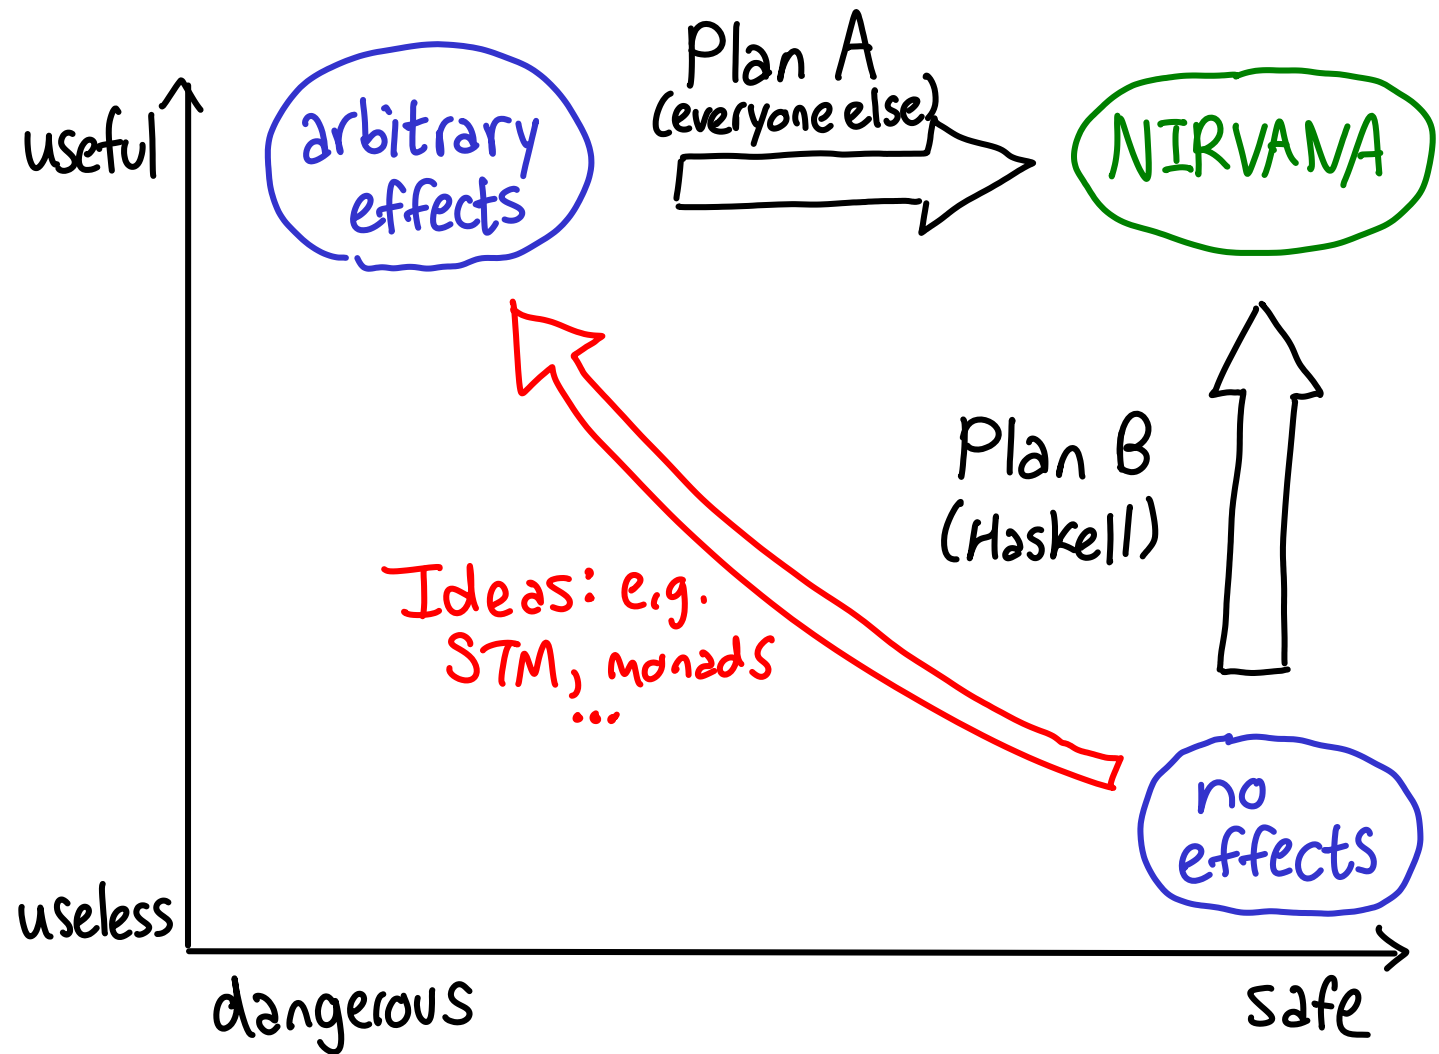
Plan: selectively permit effects

Examples:

- Domain specific languages (SQL, XQuery, MapReduce)
- Functional languages + controlled effects (Haskell)







One of Haskell's most significant contributions is to take purity seriously and relentlessly pursue Plan B.

Imperative languages will embody growing (and checkable) pure subsets.

— Simon Peyton Jones

Conclusion

- Atomic blocks raise the level of abstraction for concurrent programming.
not assembly!
- Not a silver bullet
bugs! concurrency is hard! fairness
- Performance hit, but it seems acceptable
HTM! Transactional boosting!