

Ownership

Edward Z. Yang

The plan:

- What's the problem?

Manual memory management

- C++ Unique pointers and ownership

Rules of ownership and borrowing

- Rust Statically checked ownership

Last time...

infinite memory

refcounting / garbage collection

finite memory

Last time...

but there are tradeoffs!
atomic incref/decref, pauses

infinite memory



refcounting / garbage collection

finite memory


**I DON'T ALWAYS
MANUALLY MANAGE MEMORY**



What is manual memory management?

Memory management on the stack:

```
void f() {  
    int x[20];
```

 array is available


```
}
```

X when scope exits, array
is freed


What is manual memory management?

Memory management on the stack:


```
void f() {  
    int x[20];
```


 array is available

```
}
```

 what if I want the array to live on?

Memory management on the heap

In C: `int* p = (int*) malloc(sizeof(int) * 4);`
  `// do some stuff to p`
 `free(p);`

In C++: `A* p = new A();`
  `// do some stuff to p`
 `delete p;`

Memory management on the heap

```
A* p = foobar();
```

```
// do some stuff with p
```

```
ok... now what?
```

Memory management on the heap


```
inline void
THPUtils_packInt64Array(
    PyObject *tuple,
    size_t size,
    const int64_t *sizes
) {
    for (size_t i = 0; i != size; ++i) {
        PyObject *i64 = THPUtils_packInt64(sizes[i]);
        if (!i64) {
            throw python_error();
        }
        PyTuple_SET_ITEM(tuple, i,
                          THPUtils_packInt64(sizes[i]));
    }
}
```

spot the bug...

Memory management on the heap

```
inline void
THPUtils_packInt64Array(
    PyObject *tuple,
    size_t size,
    const int64_t *sizes
) {
    for (size_t i = 0; i != size; ++i) {
        PyObject *i64 = THPUtils_packInt64(sizes[i]);
        if (!i64) {
            throw python_error();
        }
        PyTuple_SET_ITEM(tuple, i,
                         THPUtils_packInt64(sizes[i]));
    }
}
```

what does this
function do?



spot the bug...

Memory management on the heap

```
inline void
THPUtils_packInt64Array(
    PyObject *tuple,
    size_t size,
    const int64_t *sizes
) {
    for (size_t i = 0; i != size; ++i) {
        PyObject *i64 = THPUtils_packInt64(sizes[i]);
        if (!i64) {
            throw python_error();
        }
        PyTuple_SET_ITEM(tuple, i,
            THPUtils_packInt64(sizes[i]));
    }
}
```

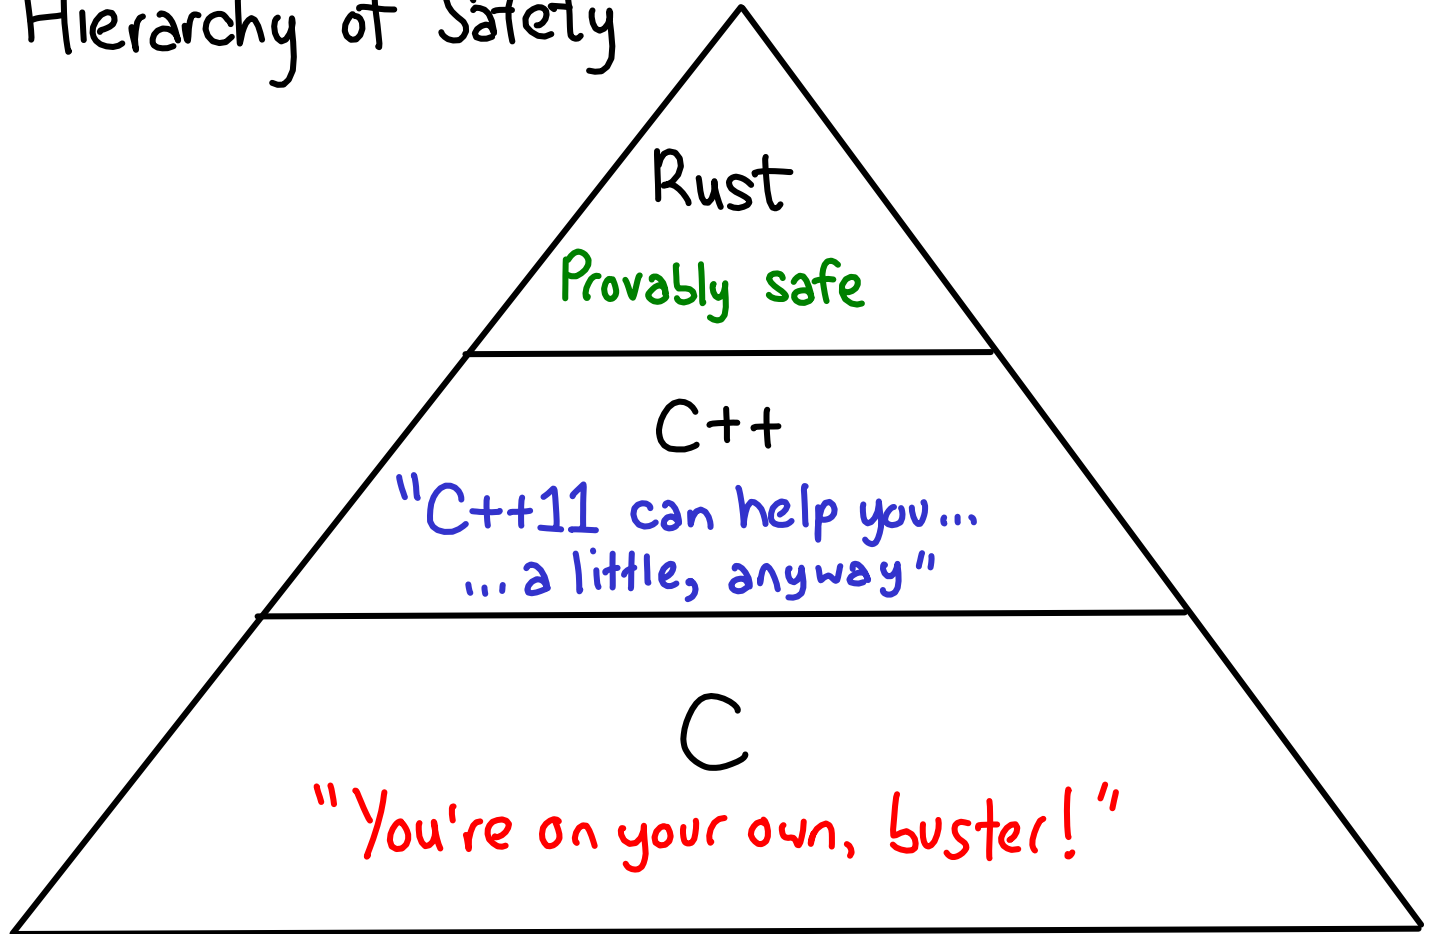
what about this one?

spot the bug...

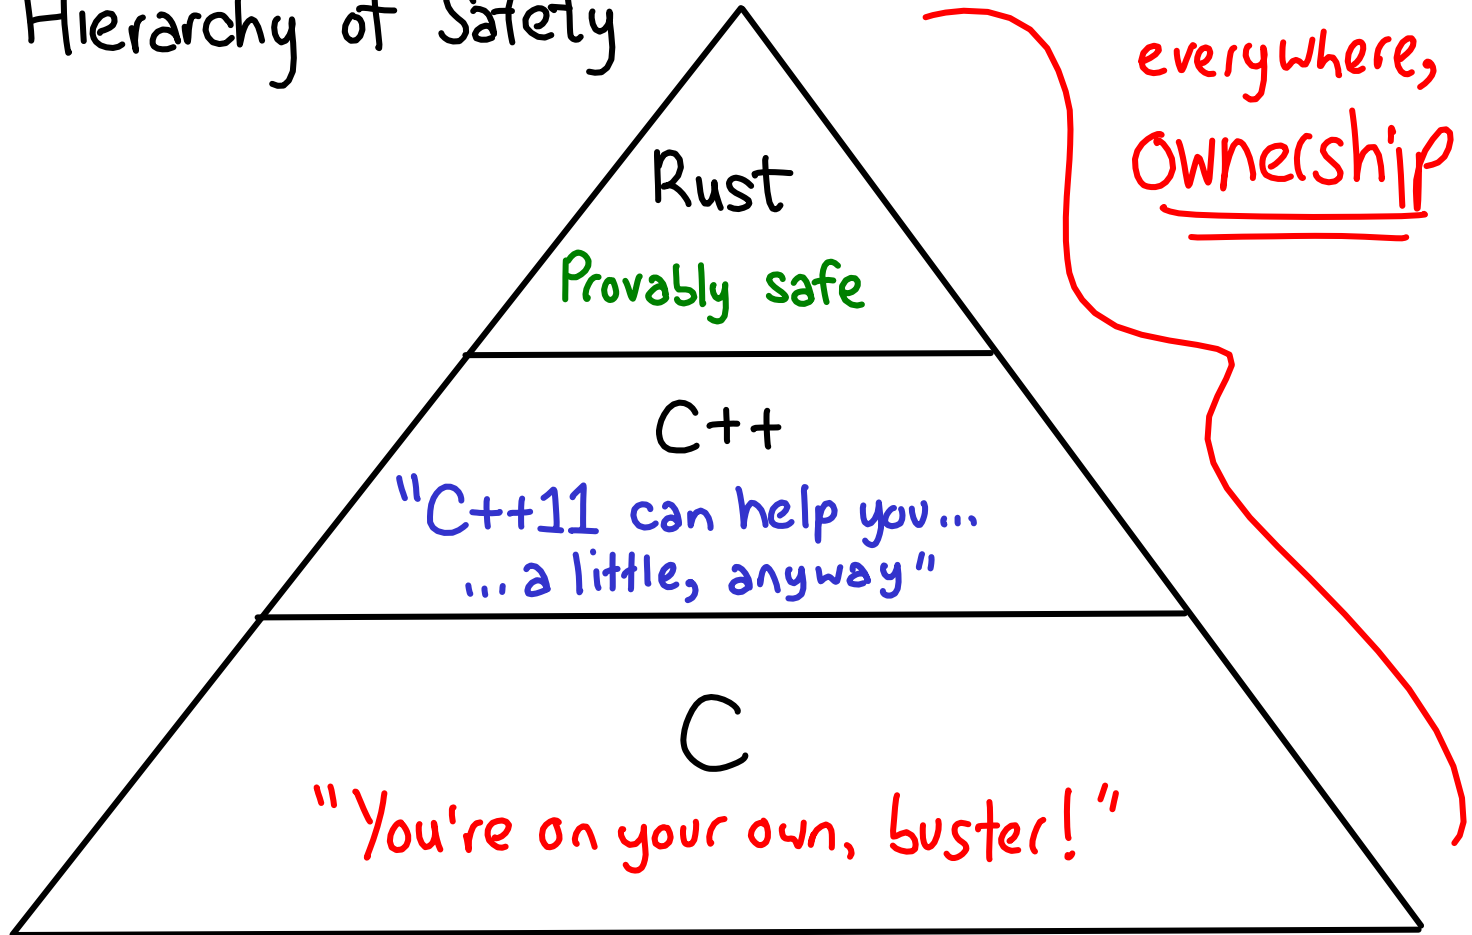
Bug farm

- I have a pointer: should I destroy it when I'm done using it?
- How should I destroy a pointer when I'm done with it?
- Did I delete exactly once in all codepaths?
- Did I already free this pointer?
(Is it dangling?)

Hierarchy of Safety



Hierarchy of Safety



What is ownership? It's a set of rules

Every allocation has a unique owner

The owner is responsible for deallocating memory when it is done using it.


```
inline void  
THPUtils_packInt64Array(  
    PyObject *tuple,  
    size_t size,  
    const int64_t *sizes
```

```
) {  
    for (size_t i = 0; i != size; ++i) {  
        PyObject *i64 = THPUtils_packInt64(sizes[i]);  
        if (!i64) {  
            throw python_error();  
        }  
        PyTuple_SET_ITEM(tuple, i,  
            THPUtils_packInt64(sizes[i]));  
    }  
}
```

allocates a new
PyObject, and transfers
ownership to the caller

← error is failing to
free i64

← steals ownership, absolving caller
of responsibility to free

Ownership is everywhere

(P.S. PyObject is refcounted: still has notion of ownership for incref/decref)

Changed in version 2.5: This function returned an `int` type. This might require changes in your code for properly supporting 64-bit systems.

`PyObject* PyTuple_GetItem(PyObject *p, Py_ssize_t pos)`

Return value: Borrowed reference.

Return the object at position `pos` in the tuple pointed to by `p`. If `pos` is out of range, return `NULL` and sets an `IndexError` exception.

Changed in version 2.5: This function used an `int` type for `pos`. This might require changes in your code for properly supporting 64-bit systems.

`PyObject* PyTuple_GET_ITEM(PyObject *p, Py_ssize_t pos)`

Return value: Borrowed reference.

Like `PyTuple_GetItem()`, but does no checking of its arguments.

Changed in version 2.5: This function used an `int` type for `pos`. This might require changes in your code for properly supporting 64-bit systems.

`PyObject* PyTuple_GetSlice(PyObject *p, Py_ssize_t low, Py_ssize_t high)`

Return value: New reference.

Take a slice of the tuple pointed to by `p` from `low` to `high` and return it as a new tuple.

Changed in version 2.5: This function used an `int` type for `low` and `high`. This might require changes in your code for properly supporting 64-bit systems.

`int PyTuple_SetItem(PyObject *p, Py_ssize_t pos, PyObject *o)`

Insert a reference to object `o` at position `pos` of the tuple pointed to by `p`.

```
PyObject *p = alloc_obj();
```

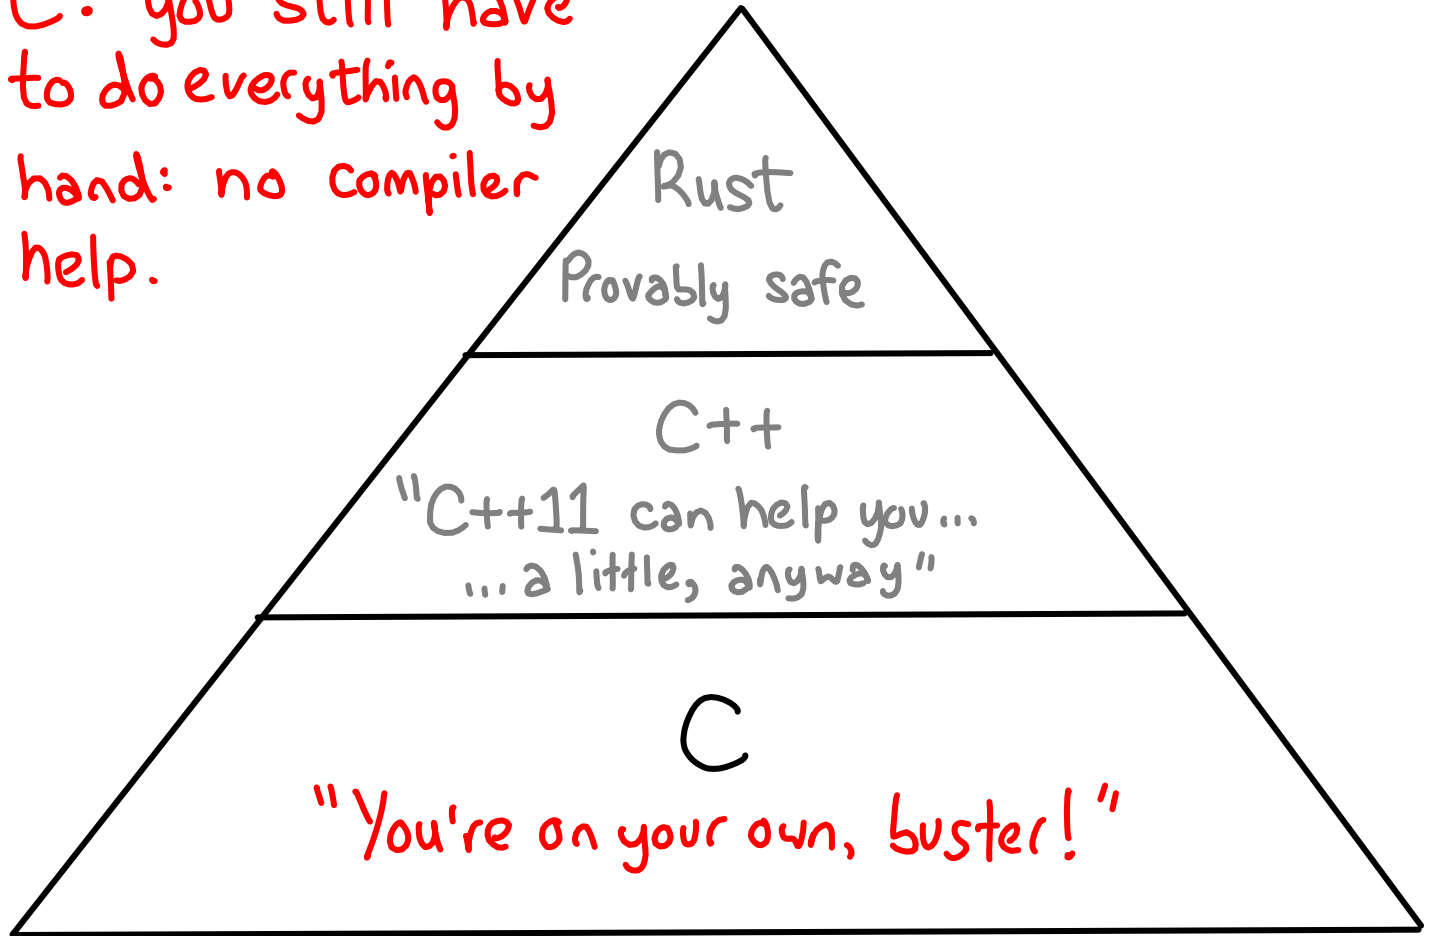
```
borrow_obj(p);
```

```
Py_CLEAR(p);
```



borrow
(function is not
responsible for
freeing)

C: you still have
to do everything by
hand: no compiler
help.



C++11 to the rescue: smart pointers

`std::shared_ptr<T>`

`std::weak_ptr<T>`

owning pointers

`const T&`

`T*`

non-owning
references

C

```
T* allocate();  
void borrow(T*);  
void free(T*);
```

Basics

```
T* p = allocate();  
borrow(p);  
free(p);
```

↖ still own p

C++11

```
std::unique_ptr<T>  
    allocate();  
void borrow(T*);  
void free(T*);
```

Basics

```
auto p = allocate();  
borrow(p.get())
```

(automatically freed)

↑ How? `std::unique_ptr`
is itself an object allocated
on the stack! (RAII)

C

```
T* allocate();  
void steal(T*);
```

Ownership transfer

```
T* p = allocate();  
steal(p);
```

(we no longer own p ; we
must not use it, and
must not free it!)

C++11

```
std::unique_ptr<T>  
    allocate();
```

```
void steal(  
    std::unique_ptr<T>  
) ;
```

compare with:

```
void borrow(T*);
```

↑
different!

Ownership transfer

```
auto p = allocate();
```

```
steal(std::move(p));
```

```
// p is now nullptr!
```

The exact mechanism by which this works is complex: it involves implicit move construction of a new `unique_ptr`. This was not added until C++11

C++11 benefits

- Automatic disposal of owning pointers (e.g., `unique_ptr`) when they exit scope
- Type level distinction between owning and non-owning pointers

C++11: The ugly bits

```
#include <memory>
```

```
struct T {  
    int x = 0;  
};
```

```
int main() {  
    auto* p = std::make_unique<T>().get();  
    p->x = 2;  
    return 0;  
}
```

↑ is this OK?

no complaints from
the compiler...

no complaints even with warnings!



```
ezyang@sabre:~$ gcc test.cpp -std=c++14 -fsanitize=address  
ezyang@sabre:~$ ./a.out
```

```
=====
```

==10991==ERROR: AddressSanitizer: heap-use-after-free
on address 0x60200000eff0 at pc 0x00000004009cc
bp 0x7ffcdb7b1090 sp 0x7ffcdb7b1080
WRITE of size 4 at 0x60200000eff0 thread T0
#0 0x4009cb in main (/home/ezyang/a.out+0x4009cb)
#1 0x7fbf7f20282f in __libc_start_main
 (/lib/x86_64-linux-gnu/libc.so.6+0x2082f)
#2 0x400888 in _start (/home/ezyang/a.out+0x400888)

C++11: The ugly bits

Automatic disposal and type-system encoded ownership are **all** you get.

You MUST ensure that borrows don't extend beyond lifetime of owning pointer.

non-owning pointer

```
#include <memory>
```

```
struct T {  
    int x = 0;  
};
```

borrow reference to T

```
int main() {  
    auto* p = std::make_unique<T>().get();  
    p->x = 2;  
    return 0;  
}
```

temporary unique_ptr<T>, dies when the statement finishes evaluating

It can be non-obvious

```
class Foo {  
    Foo(const Foo&); // copy-constructor  
    Field* mutable_field();  
}
```

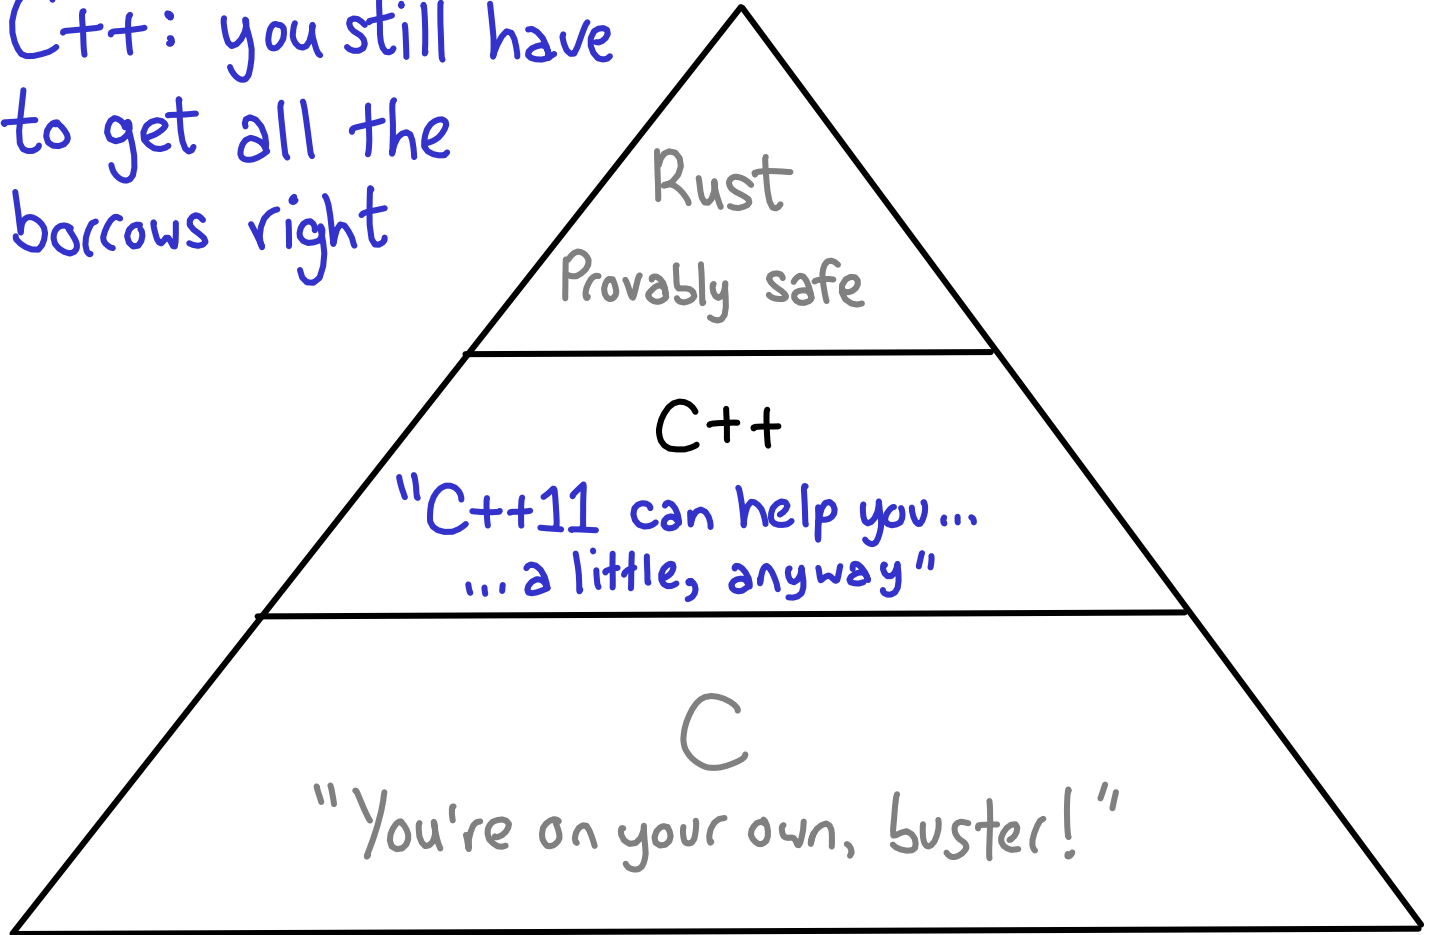
```
class Field {  
    void set_int(int);  
}
```

```
const Foo& borrow_foo();  
void run_with_borrow(Field*);
```

```
void run_with_modified_copy_of_field() {  
    auto* field = Foo(borrow_foo()).mutable_field();  
    field->set_int(23);  
    run_with_borrow(field);  
}
```

↑ trouble! temporary!

C++: you still have
to get all the
borrows right



Rust

Rust

- Conceptually, same principles as C++
 - No mutable aliases: fearless concurrency
- All borrows checked by the borrow checker
- Some programs not expressible in Safe Rust (doubly-linked list)
 - unsafe escape hatch

Rust: safe, but less
programs allowed!

